

10. Capacity Optimization

Special Topics in Computer Systems:

Modern Storage Systems

(IC820-01)

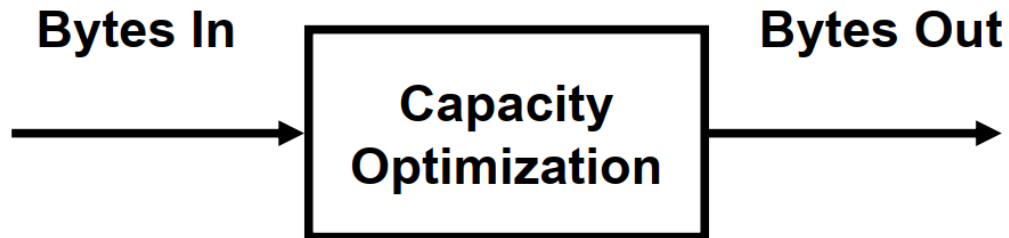
Instructor:

Prof. Sungjin Lee (sungjin.lee@dgist.ac.kr)

Capacity Optimization

- There are two popular methods to improve storage capacity
 - Method 1: *Data Deduplication*
 - The replacement of multiple copies of data with references to a shared copy in order to save storage space
 - Method 2: *Data Compression*
 - The encoding of data to reduce its storage requirements
- Widely used in modern storage systems and devices
 - Not only improve *storage capacity*, but *I/O performance* and *energy efficiency*

Space Reduction Ratio & Percent



$$\text{Ratio} = \frac{\text{Bytes In}}{\text{Bytes Out}}$$

$$\% = \frac{\text{Bytes In} - \text{Bytes Out}}{\text{Bytes In}}$$

$$\text{Ratio} = \left(\frac{1}{1 - \%} \right)$$

$$\% = 1 - \left(\frac{1}{\text{Ratio}} \right)$$

Space Reduction Ratio & Percent (Cont.)

Space Reduction Ratio	Space Reduction Percent
2:1	$1/2 = 50\%$
5:1	$4/5 = 80\%$
10:1	$9/10 = 90\%$
20:1	$19/20 = 95\%$
100:1	$99/100 = 99\%$
500:1	$499/500 = 99.8\%$



- Ratios can meaningfully be compared only under the same set of assumptions
- Relatively low space reduction ratios provide significant space savings

Outline

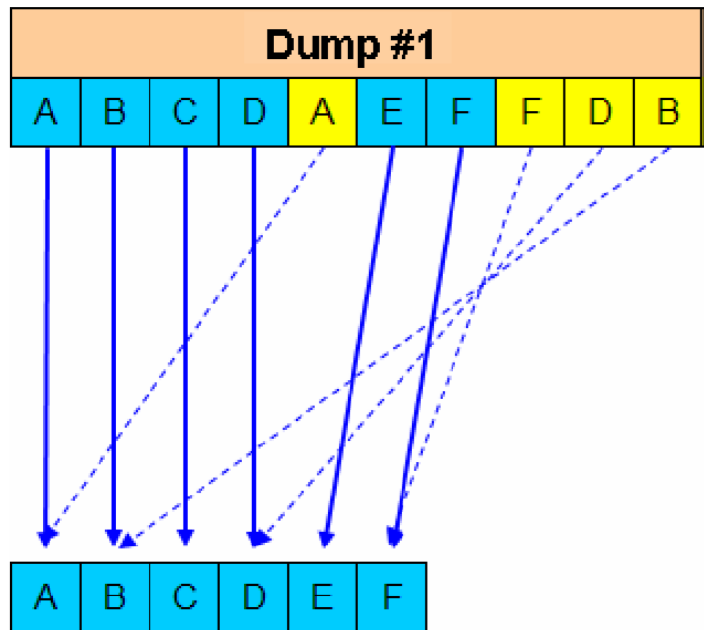
- **Data Deduplication**
- Data Compression
- Case Studies




Data Deduplication Simplified

Dump #1									
A	B	C	D	A	E	F	F	D	B


 = new unique data
 = repeat data

Data Deduplication Simplified (Cont.)

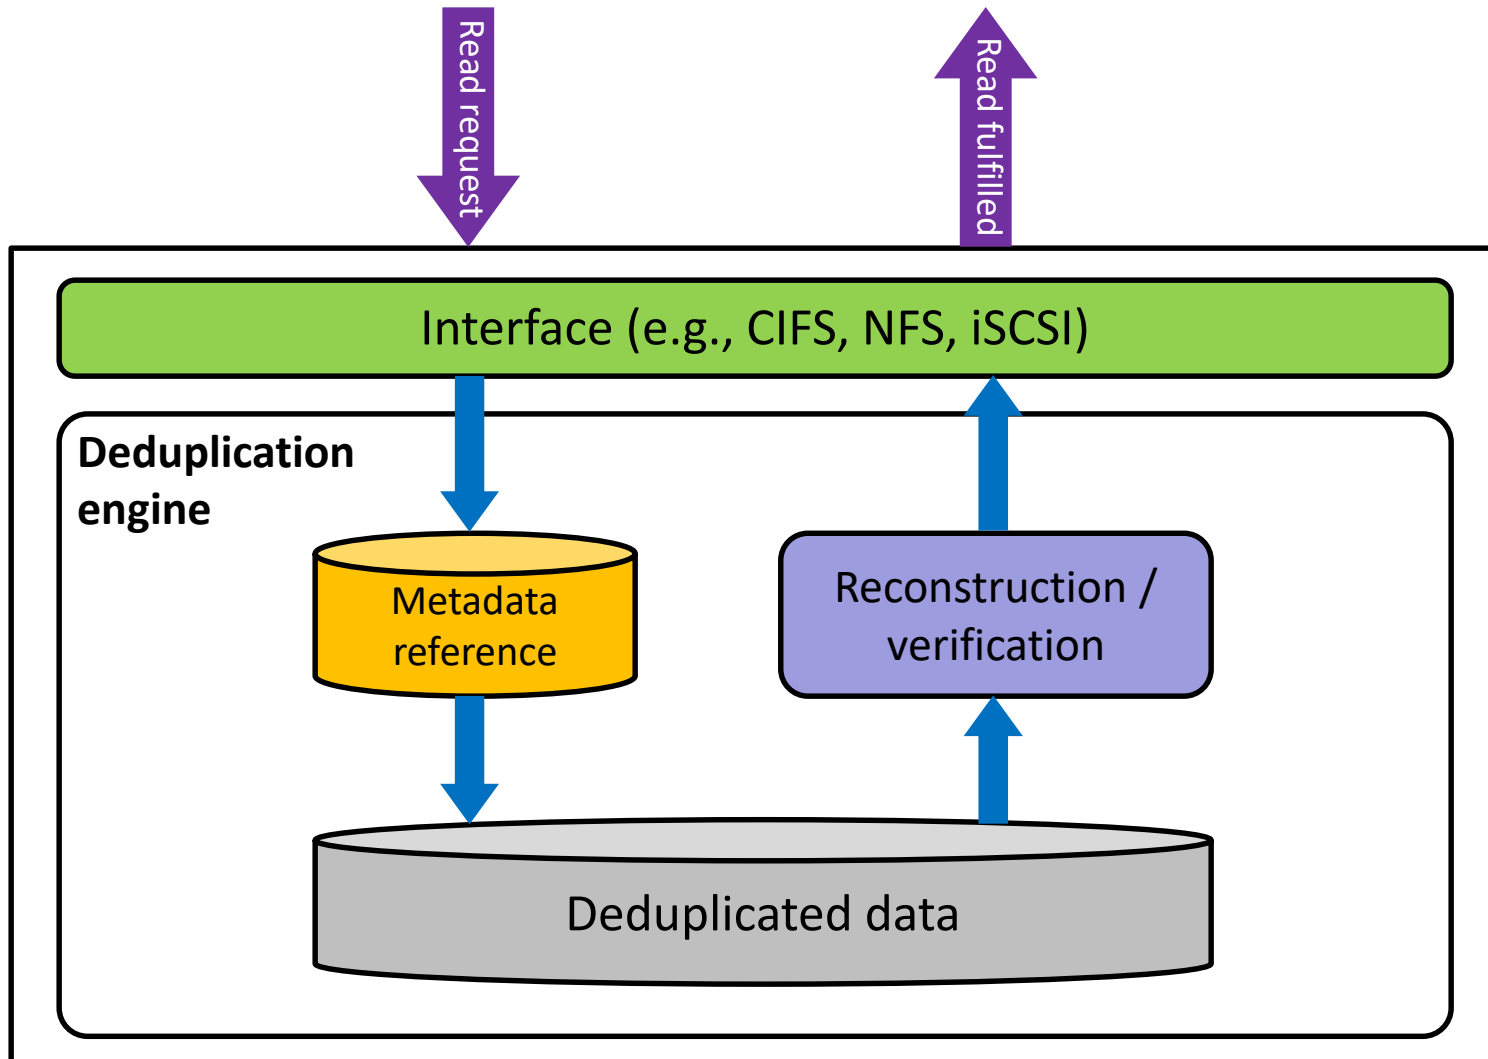


 = new unique data
 = repeat data
 = pointer to unique data segment

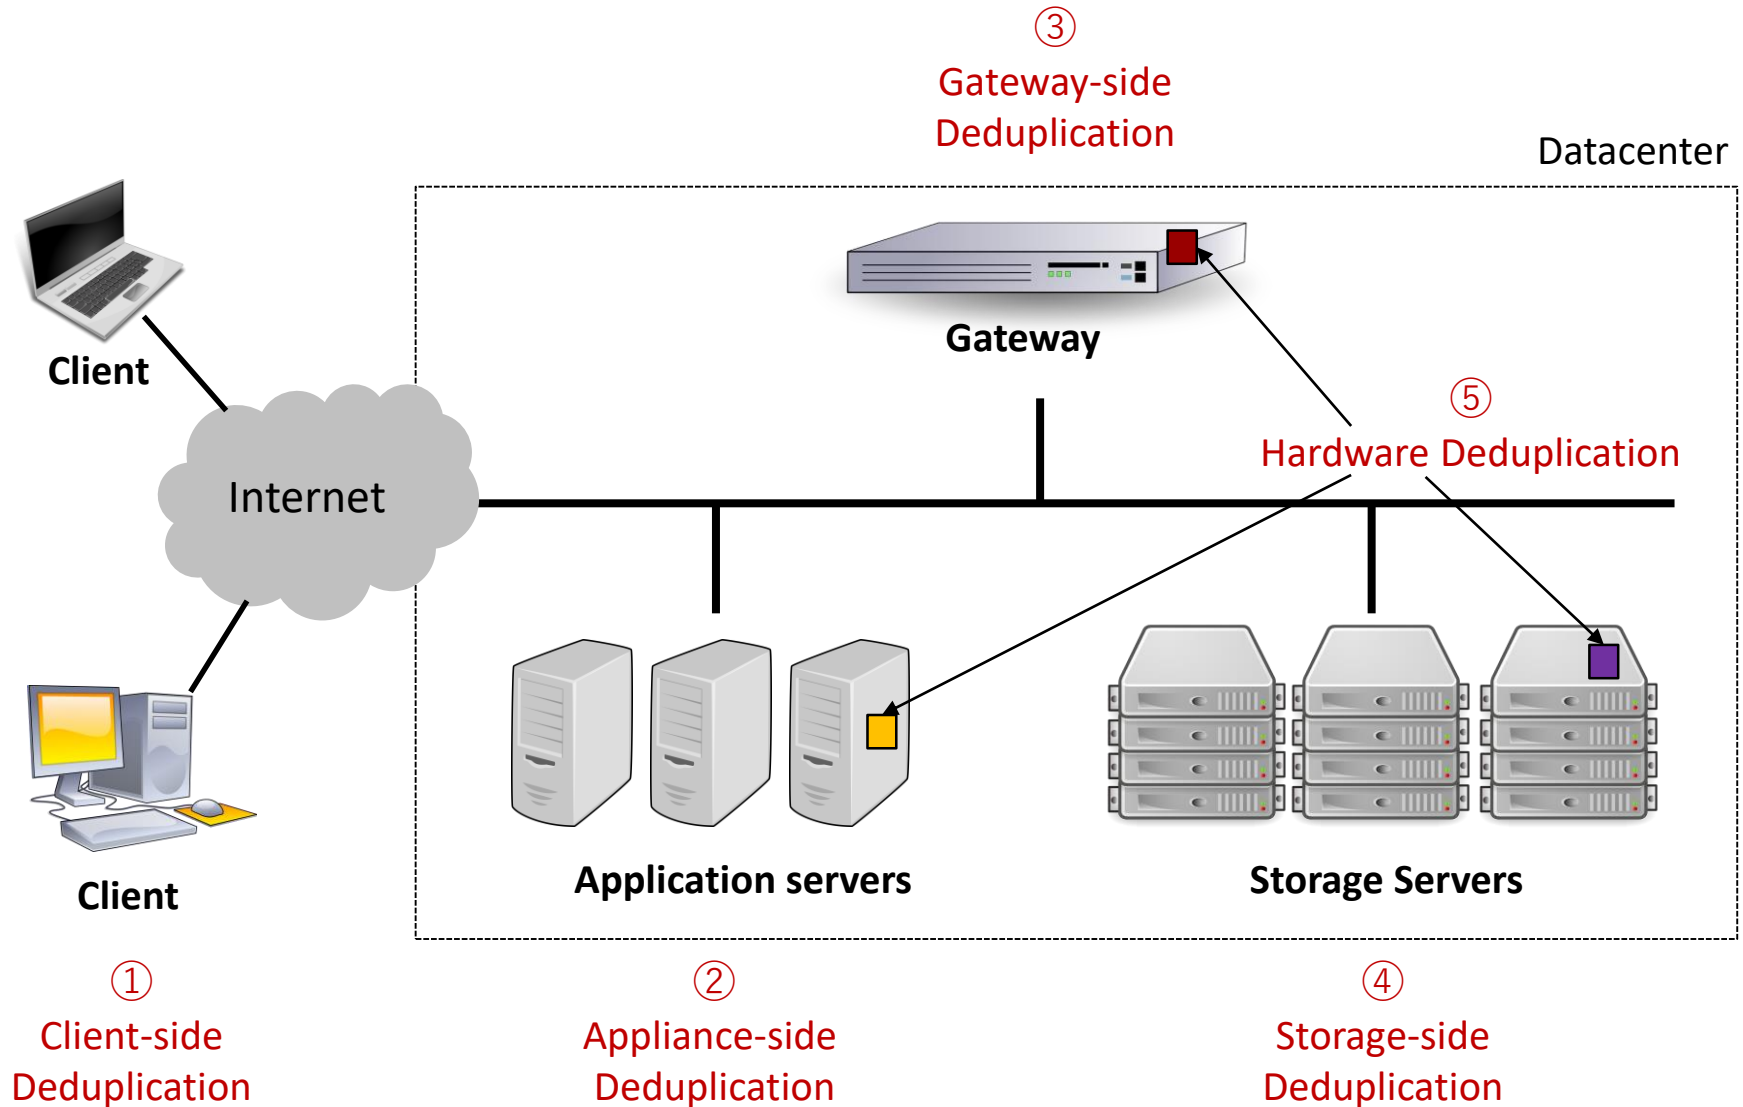
Data Deduplication Simplified (Cont.)



Reading Duplicated Data



Where is Deduplication Deployed?



Source and Target

■ *Source Deduplication*

- Identifies duplicate data at the client
- Transfers unique segments to a central repository
- Separate client and server components

■ *Target Deduplication*

- Identifies duplicate data where the data is being stored
- Stores unique segments
- Standalone systems

■ **Consideration**

- Neither approach enables a greater or lesser space savings
- Scope of data deduplication may vary by implementation

Source and Target (Cont.)

Application and Data on
Customer Premises



Deduplicated Data



Deduplication performed
by the backup client

Backup as Service
within the Cloud



Application and Data on
Customer Premises



Data



Deduplication Target
within the Cloud



Inline or Post-Process

■ *Inline Deduplication*

- Data deduplication performed *before* writing the duplicate data
- May improve I/O performance
- Low deduplication ratio

■ *Post-process Deduplication*

- Data deduplication performed *after* the data to be deduplicated has been initially stored
- May degrade I/O performance
- High deduplication ratio

Fixed or Variable Size Segment

■ *Subfile Data Deduplication*

■ *Fixed-length Segment Deduplication*

- Evaluation of data includes a fixed reference window used to look at segments of data during deduplication process
- Provide fixed granularity (e.g., 4 KB, 8 KB, or 128 KB)

■ *Variable-length Segment Deduplication*

- Evaluation of data uses a variable length window to find duplicate data in stream or volume of data processed
- Provide variable granularity

■ *Single Instance Deduplication*

- Operate at a granularity of an entire file or data object

Outline

- Data Deduplication
- **Data Compression**
- Case Studies

Data Compression

■ Lossless versus lossy compression

- Lossless compression means that no information is lost when a file is compressed and then uncompressed
- Lossy compression usually results in better compression ratio, but some information is lost

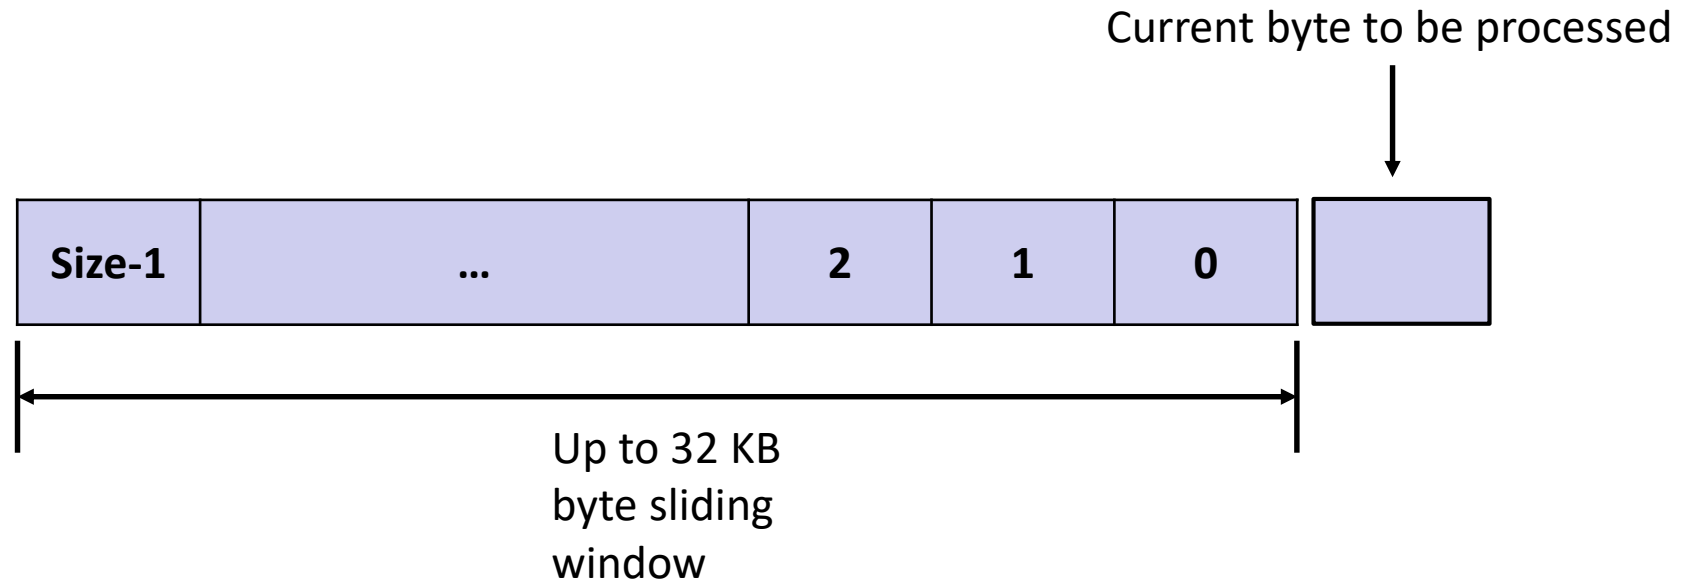
■ For data storage, lossless compression is mainly used

- LZ1-Based algorithm combined with Huffman encoding
 - LZ1 algorithm to identify matches in a sliding window history buffer
 - A post encoder to Huffman encode the matches and literals
- Hardware and software implementations

LZ1 Architecture

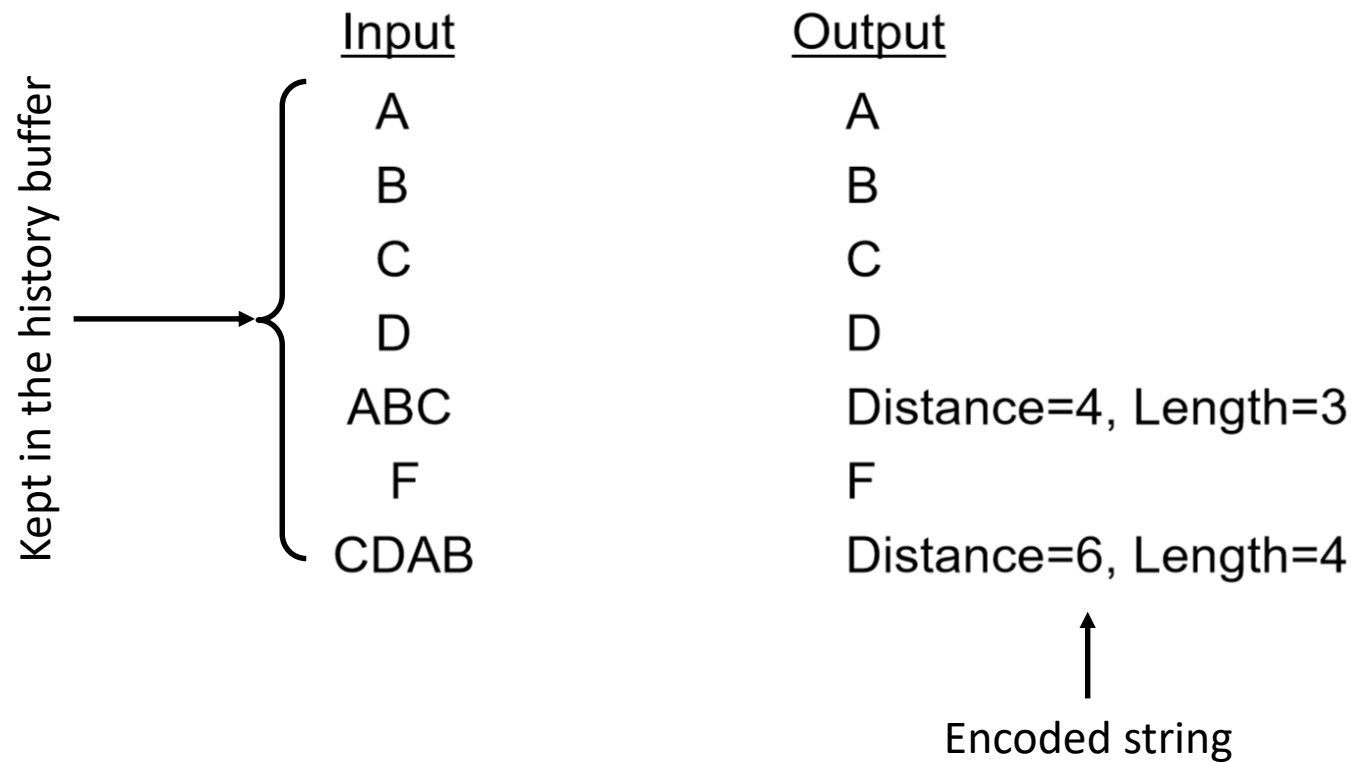
- A ***String-Matcher*** searches a History-Buffer to find repeating strings of bytes
- A sliding-window ***History-Buffer*** adds one new byte and drops off one byte from the back end of the buffer
- A ***Post-Encoder*** is a prefix encoder
 - Use statistics to encode the most common string matches with a smaller number of bits

LZ1 Algorithm and History Buffer



LZ1 String Matching

Input String: ABCDABCFCDAB.....




Huffman Encoding

Input String: A B B C B A D B

Probability of Occurrence

<u>Input character</u>	<u>Probability</u>
A	0.25
B	0.5
C	0.125
D	0.125

Huffman Encoding (Cont.)



<u>Symbol</u>	<u>Code</u>	<u>Pr</u>
A	10	0.25
B	0	0.5
C	110	0.125
D	111	0.125

- ♦ Reduction = $\frac{1}{2}[0.25(2) + 0.5(1) + 0.125(3) + .125(3)] = 0.875$
- ♦ Reduction in data size due to Huffman encoding.

Data Dependent

- Random data provides poor compression ratio performance
- Data with repeating byte strings, 2-byte or longer provides grater compression ratio performance
- Compression ratio greater than 100:1 are possible
- May expand if attempting to compress previously compressed data, but a system could detect this and send original data without compression

Algorithm Dependent


- Size of sliding window
- Static or dynamic Huffman encoding
- Number of matches tracked
- Length of matches the algorithm will search for

Hardware Implementation

- **Software compression may result in the degradation of I/O throughput and latency as well as energy consumption**
- **Hardware compression**
 - (+) Higher data rate throughput (10x)
 - (+) Free up valuable CPU bandwidth
 - (+) reduce power consumption
 - (+) Speed up a network link by sending shorter files
 - (–) Lower compression ratio
- **Several hardware-specific compression algorithms are available**
 - Hardware implementation of GZIP algorithms
 - X-Match PRO: Real-time compression / decompression

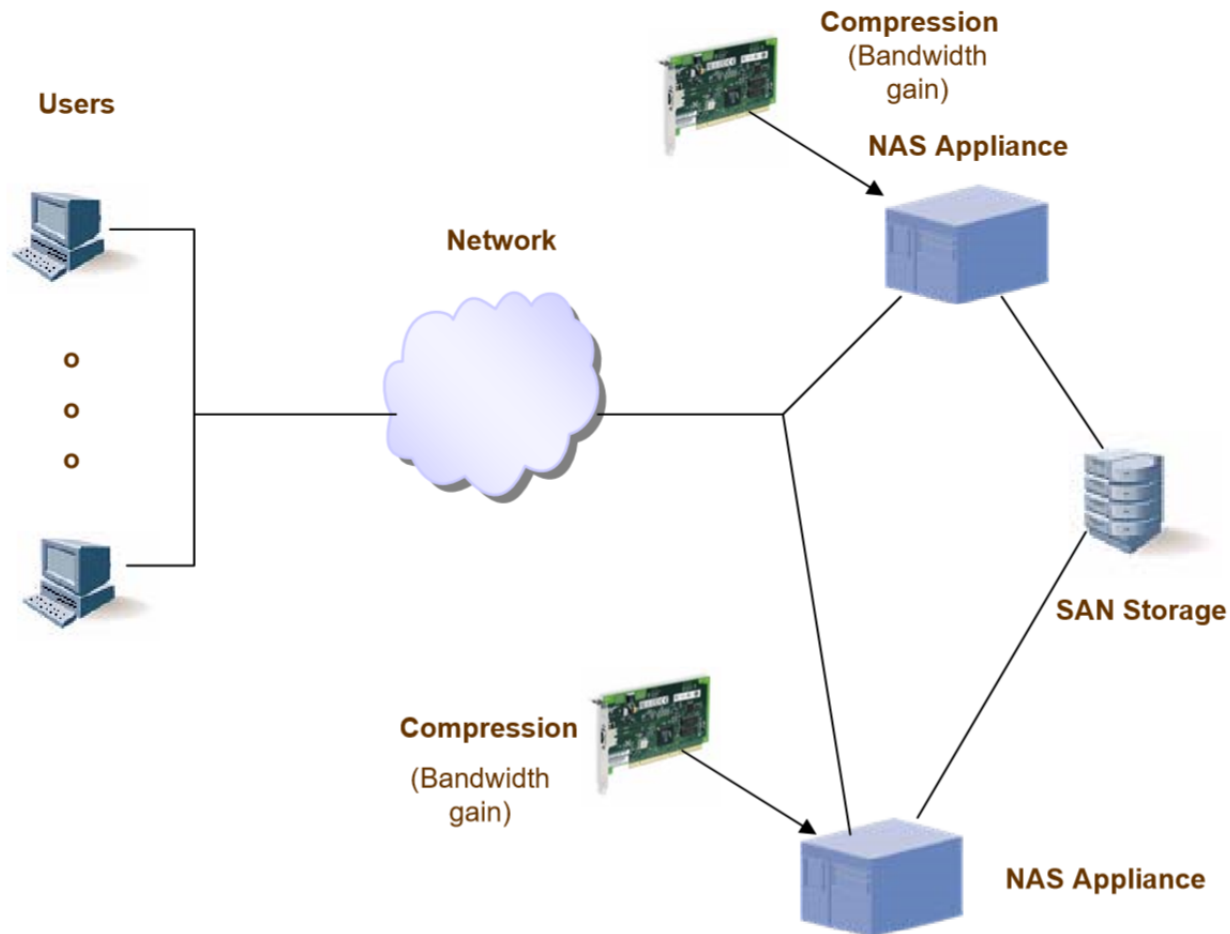
Deployment of Compression

- Similar to deduplication; combined with deduplication



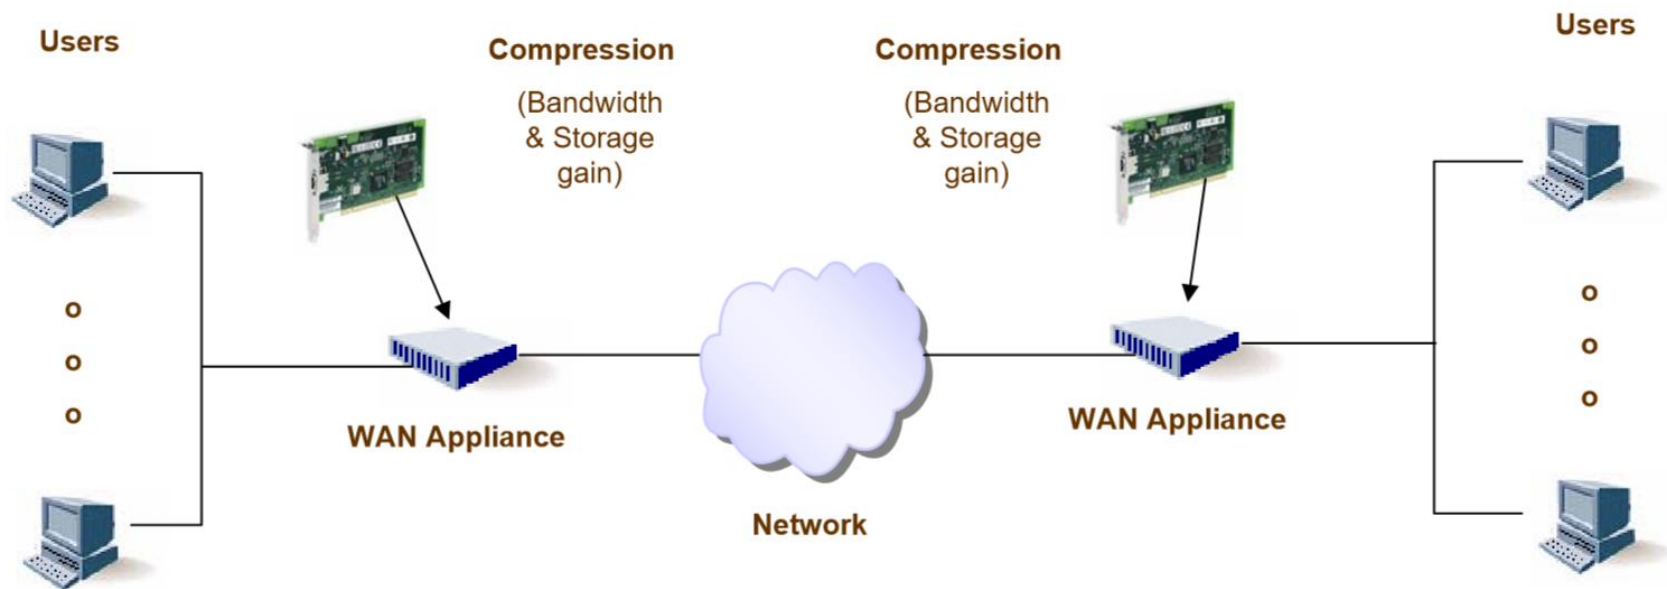
Deployment of Compression (Cont.)

- Similar to deduplication; combined with deduplication



Deployment of Compression

- Similar to deduplication; combined with deduplication



Data Deduplication vs Data Compression

■ Data deduplication

- (+) Large space saving is possible for exactly matched segments or files
- (–) Even a single-bit difference makes deduplication ineffective
- (–) Unable to get rid of repeated bit patterns of segments or files

■ Data compression

- (+) Effectively remove repeated bit patterns from input data streams
- (–) Do not provide a high compression ratio all the times

■ Common design practice

- Use deduplication for input data streams
- Use lossless compression for storage of duplicate data and remaining meta data

Outline


- Data Deduplication
- Data Compression
- **Case Studies**
 - **CA-SSD: Content-aware Solid-state Drives**
 - BlueZIP: Hardware-accelerated Compression
 - DAC: Dedup-assisted Compression

Value Locality


- To improve the reliability and performance of SSD, many FTL schemes exploit temporal/spatial locality
 - Temporal locality: buffering writes to eliminate duplicate writes
 - Spatial locality: coalescing multiple sub page writes into fewer page writes
- However, another form of locality, *Value Locality*, has not been completely unexplored
 - Value locality: certain content is accessed preferentially

Value Popularity (VP)


- Value popularity: the number of occurrences of each unique value
- Value popularity in real-world workloads



(a) web



(b) mail



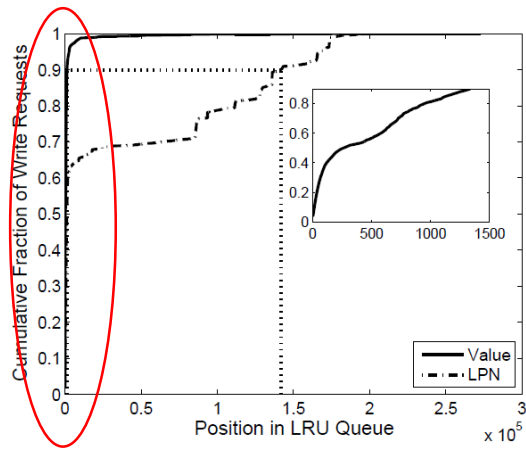
(c) homes

■ Workload statistics

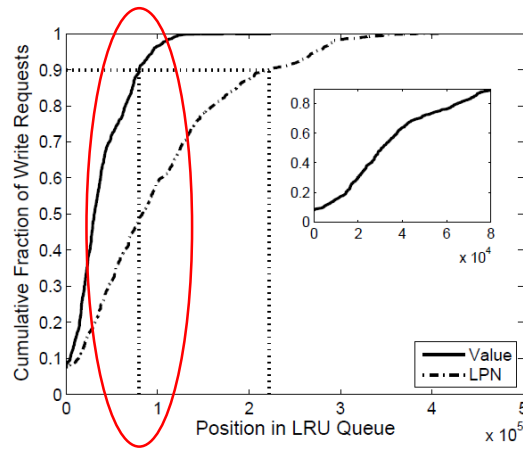
Workload	Size (GB)	% Writes	Req. (mill.)	Unique Request (%)		Seq. %
				Write	Read	
web	1.95	77.01	3.8	42.35	32.05	83.8
mail	4.22	77.32	3.6	7.83	80.85	94.7
homes	3.02	96.76	4.4	66.37	80.75	70.8

Temporal Value Locality (TVL)

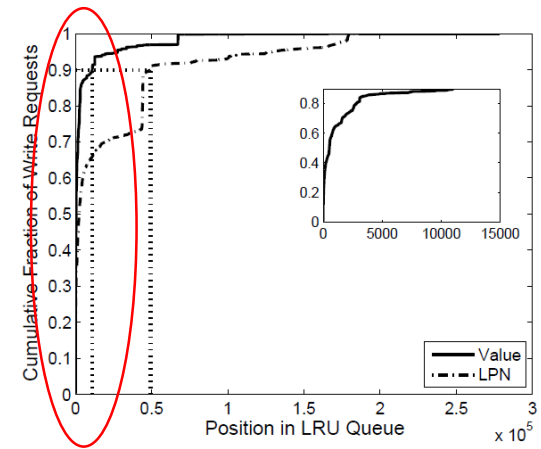
- **Temporal Value Locality:** If a certain value is accessed now, it is likely to be accessed again in the near future
- **Temporal value locality in real-world workloads**
 - The metadata cache is managed as a queue with an LRU eviction policy
 - CDFs of number of writes of the value at the $(i+1)$ st location in the LRU queue



(a) web




(b) mail



(c) homes

Cache Miss Rate for Popular Values

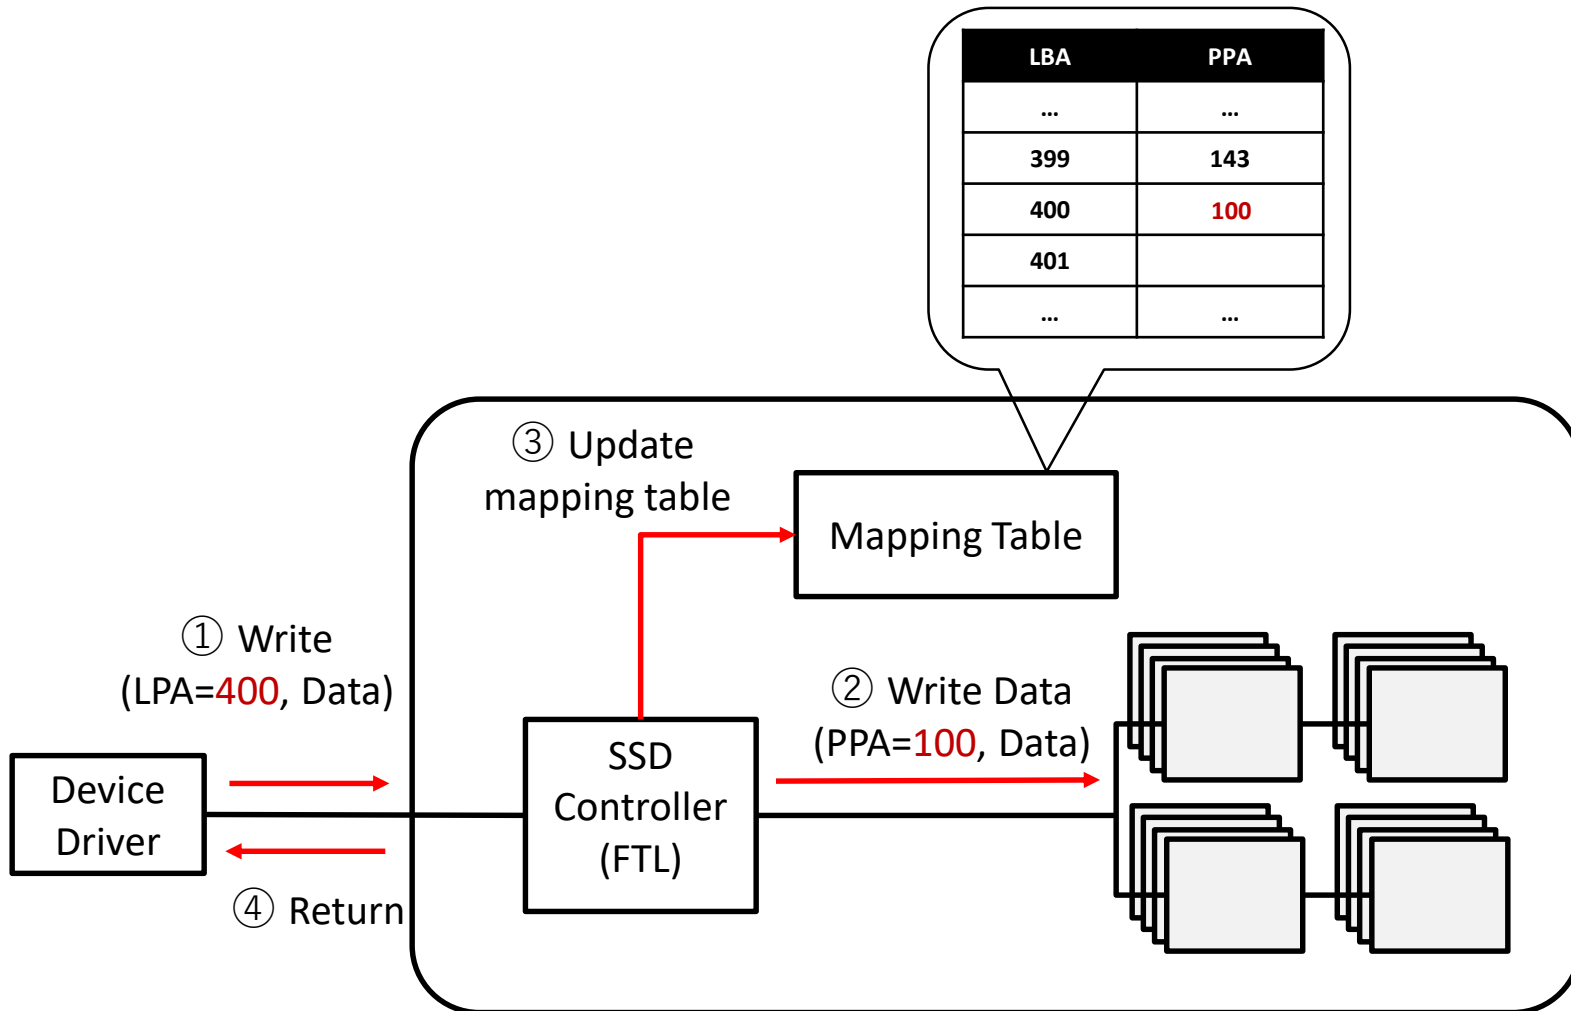


- Popular values represent the minimum number of values which account for 50% of accesses
- All the three traces achieve about 90% hit rate with 1.75 MB metadata cache (64 K hashes X 28 B = 1.75 MB)

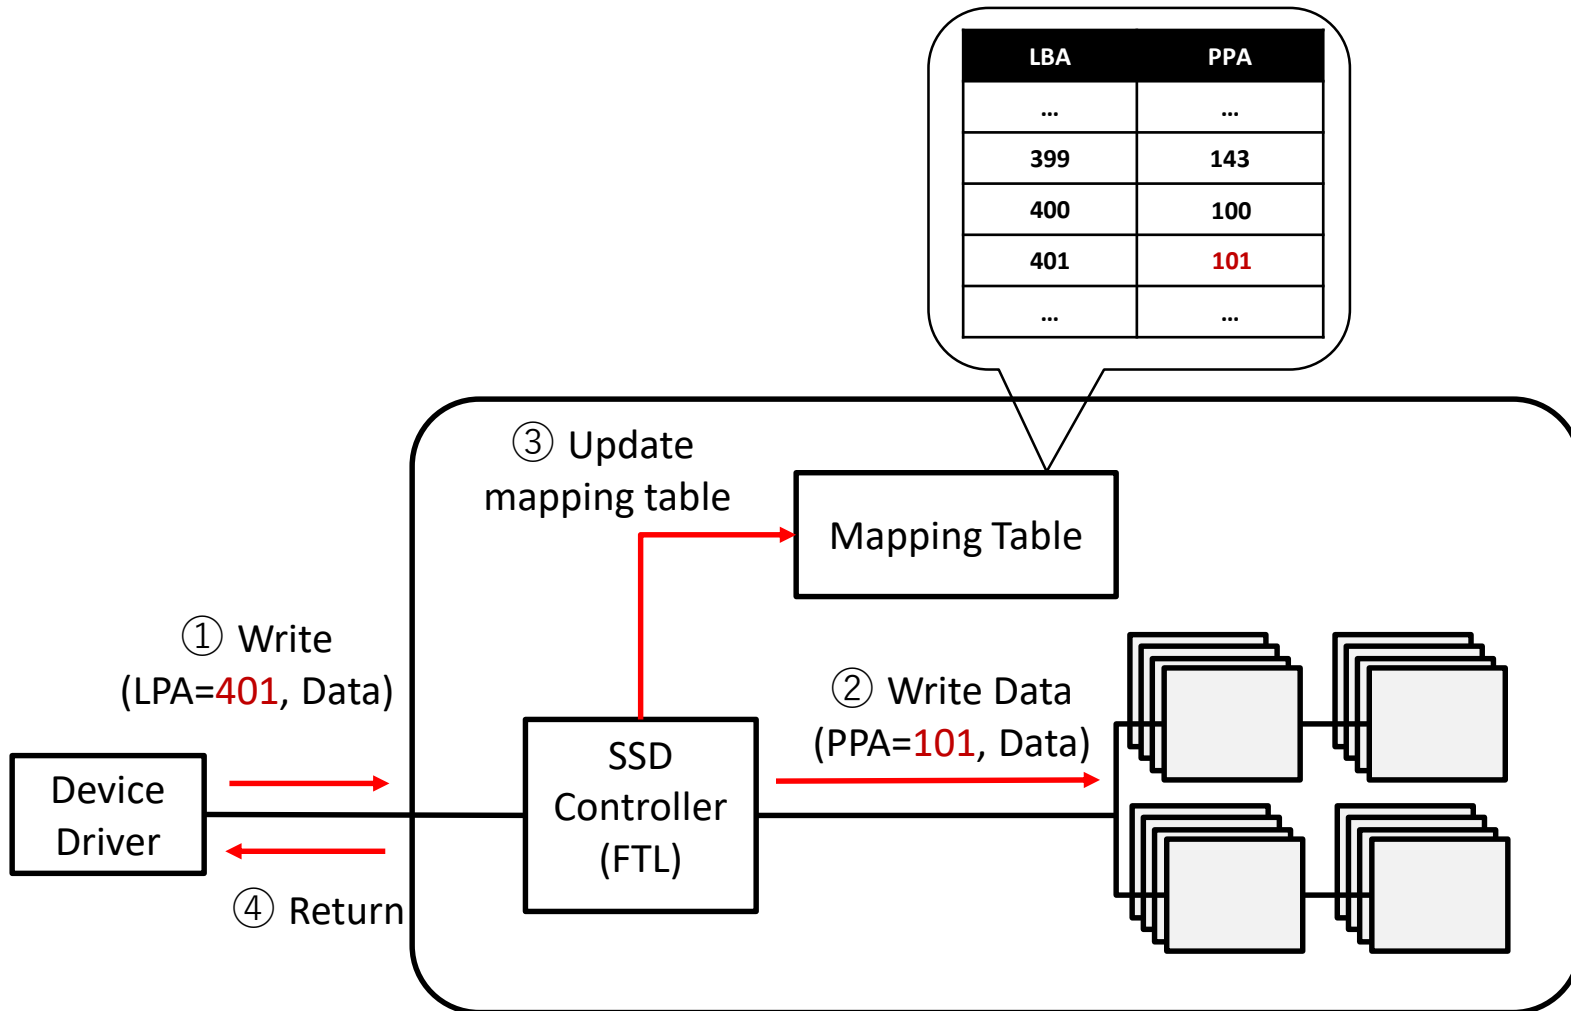
Content-Aware SSD

- The presence of value locality in a workload means that it preferentially accesses certain content over others
- This property facilitates data de-duplication inside an SSD
 - Store only a non-intersecting chunk having *a unique hash value*
 - Other data blocks with the same content point to the unique block

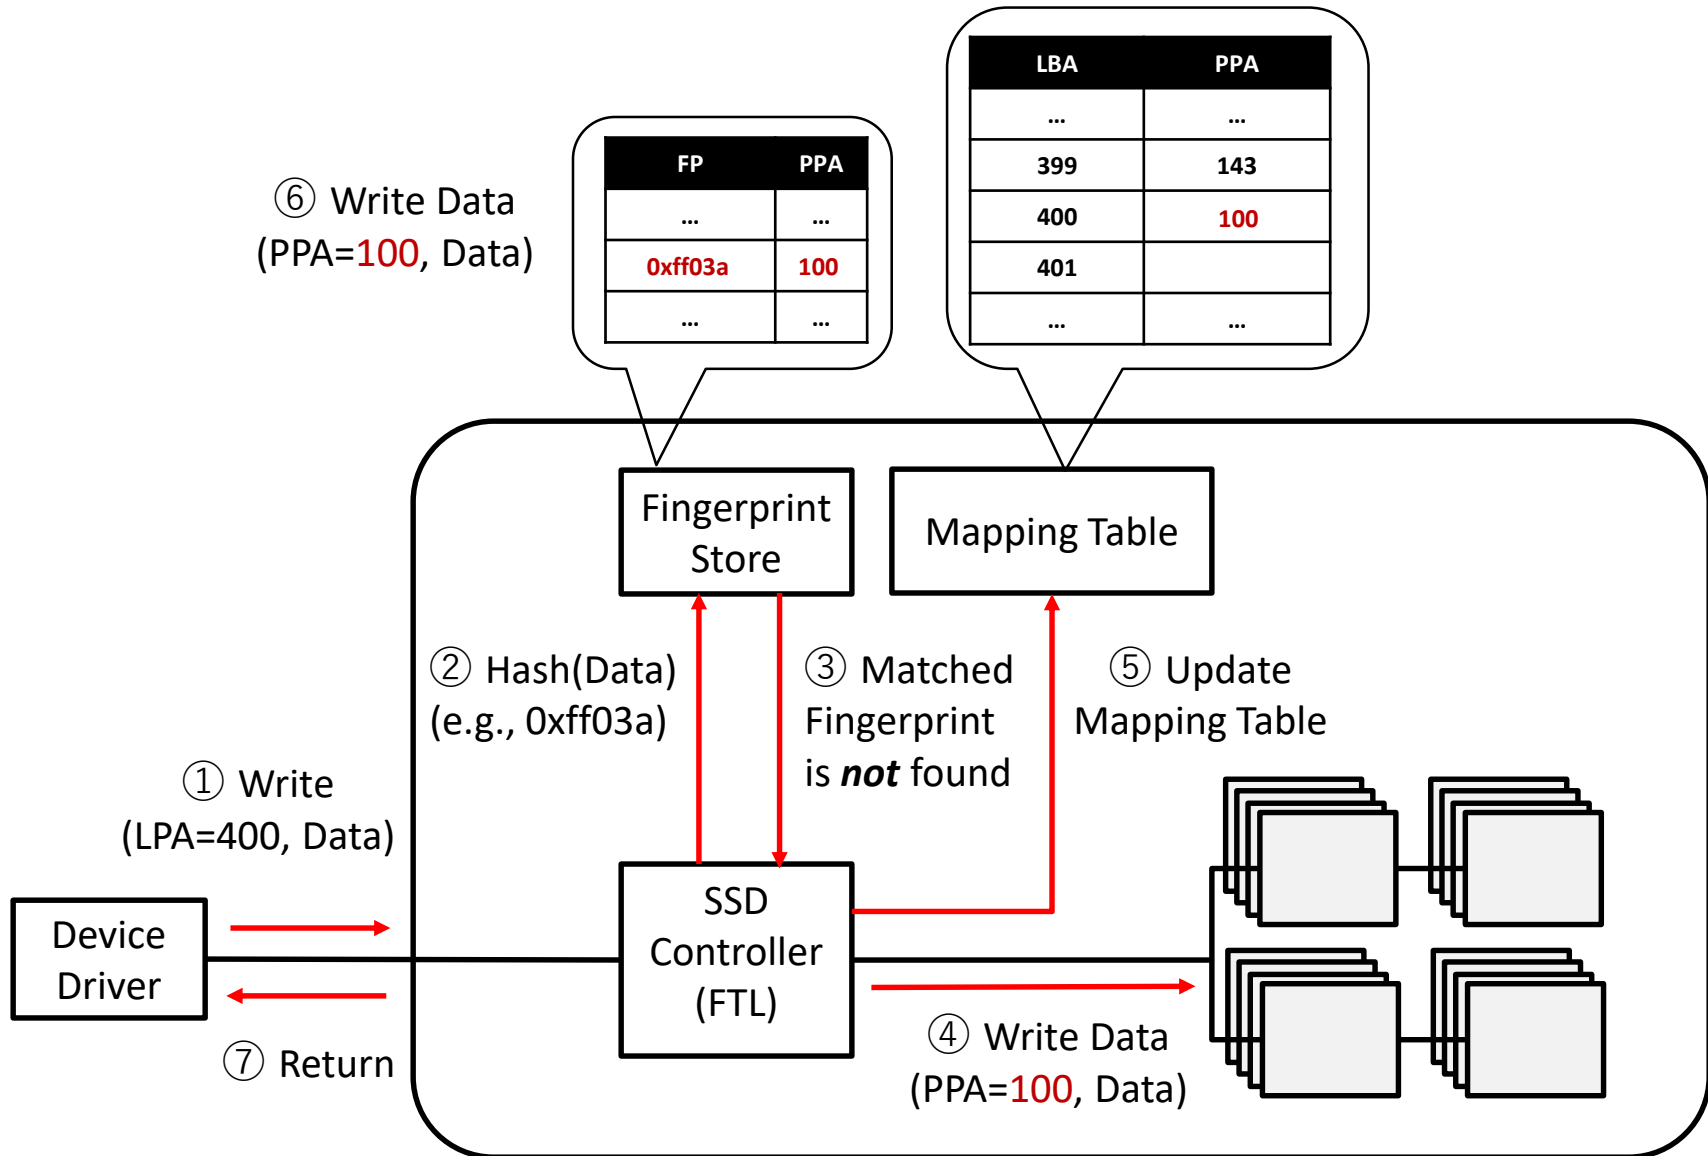
How a Conventional SSD Works?



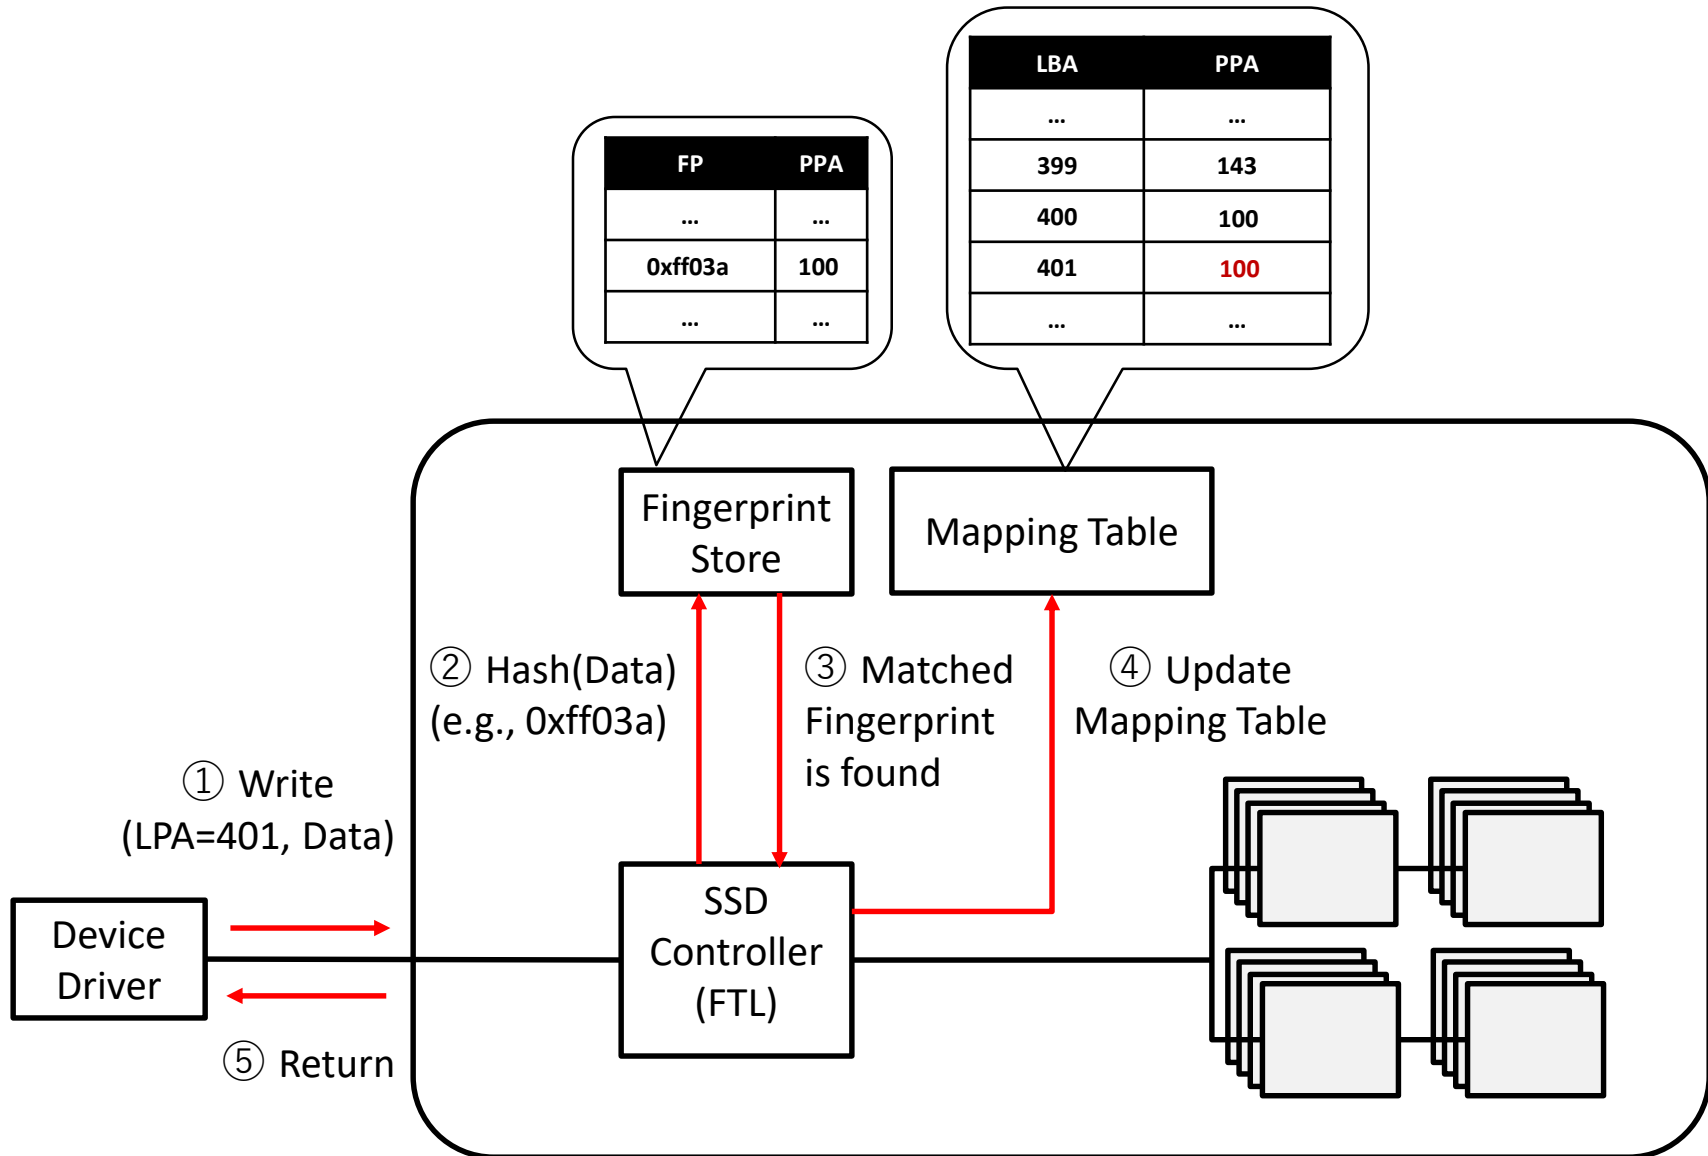
How a Conventional SSD Works?



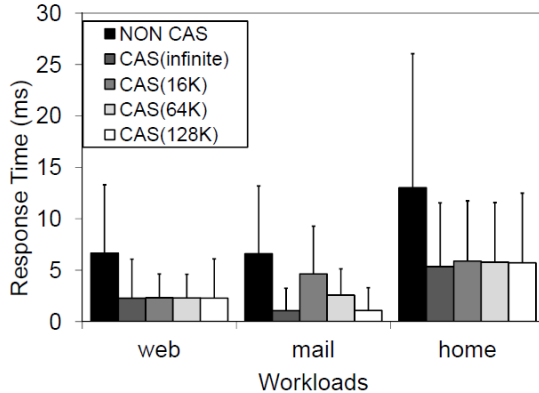
How CA-SSD Work? (Simplified)



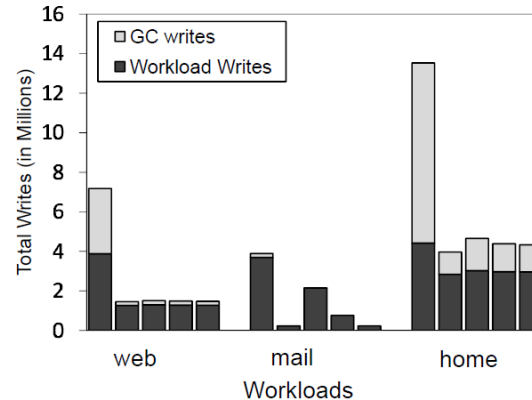
How CA-SSD Work? (Simplified) (Cont.)



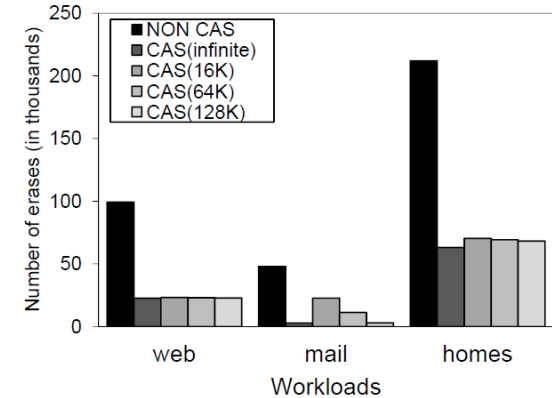
Experimental Results



(a) Response Time



(b) Total Writes



(c) Block Erases

- The reduction in write traffic: 77%, 93%, and 70% for web, mail, and home
- The write reduction benefits directly translate into the reduced response time and the reduced block erases
- CAS(128K) provides the performance close to the CA(infinite)
 - mail shows lower TVL and requires a larger metadata cache

Outline

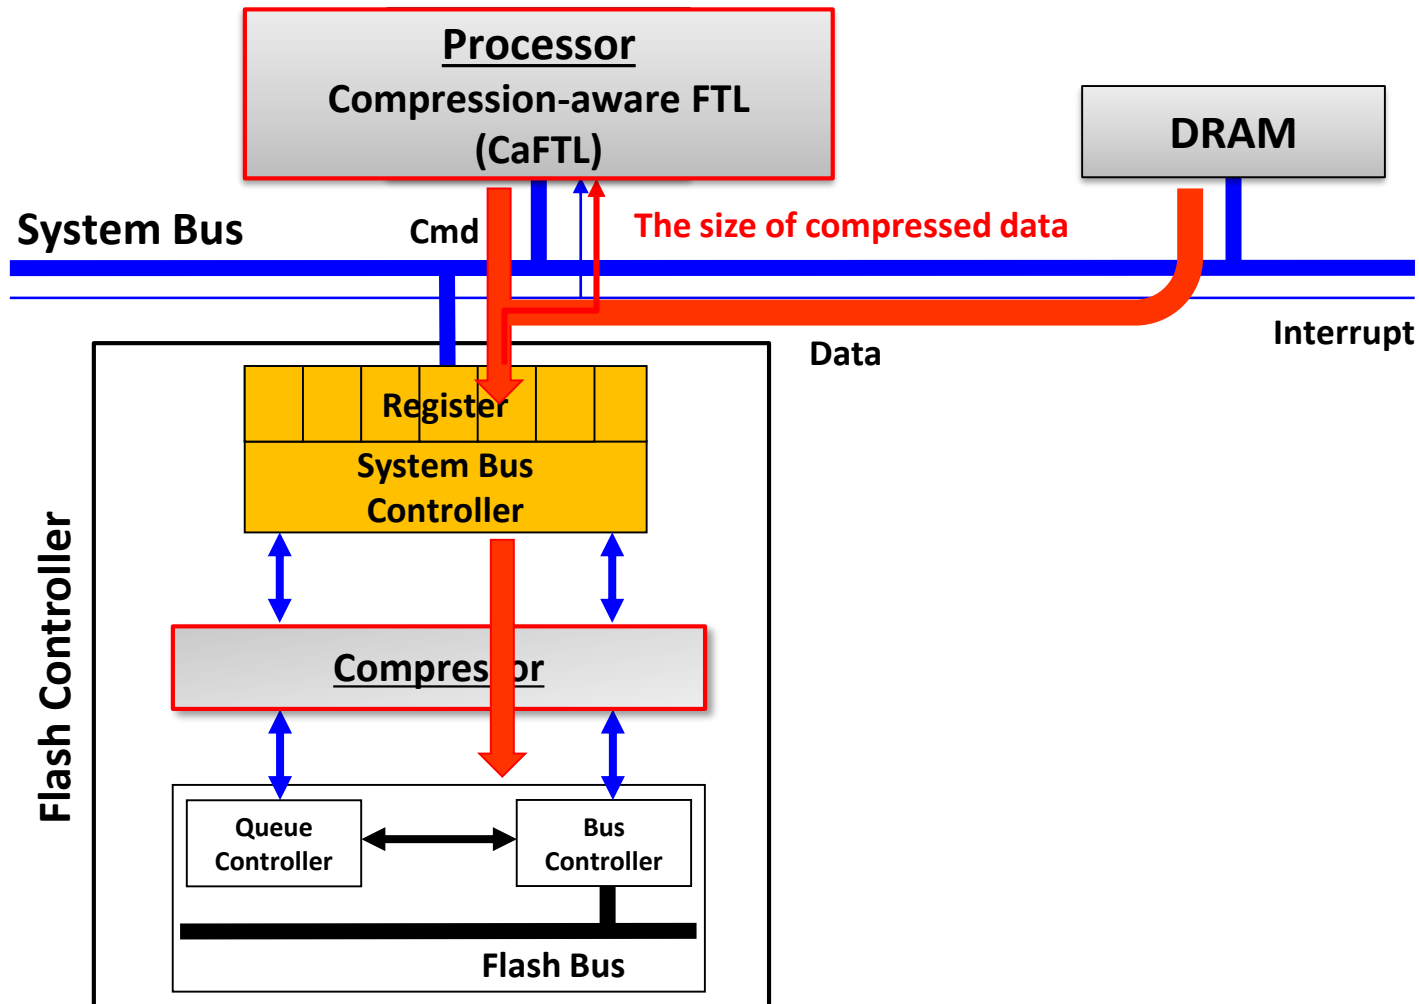
- Data Deduplication
- Data Compression
- **Case Studies**
 - CA-SSD: Content-aware Solid-state Drives
 - **BlueZIP: Hardware-accelerated Compression**
 - DAC: Dedup-assisted Compression

BlueZIP: Hardware-accelerated Compression

- **Requested data contain lots of repeated bit patterns**
 - Lossless compression helps to eliminate such repeated bit patterns.
- **BlueZIP: A hardware-accelerated compression technique,**
 - Compress requested data at runtime
 - Use a hardware-accelerated compression module
 - Provide software support for maximizing the benefits of hardware-accelerated compression
 - Improve the lifetime and performance of storage devices
 - Reduce the amount of data written to flash memory
 - Reduce the time taken to transfer data between a host system and flash memory

Overall Architecture of BlueZIP

- Implement the hardware compression module inside the flash controller to reduce system bus traffic



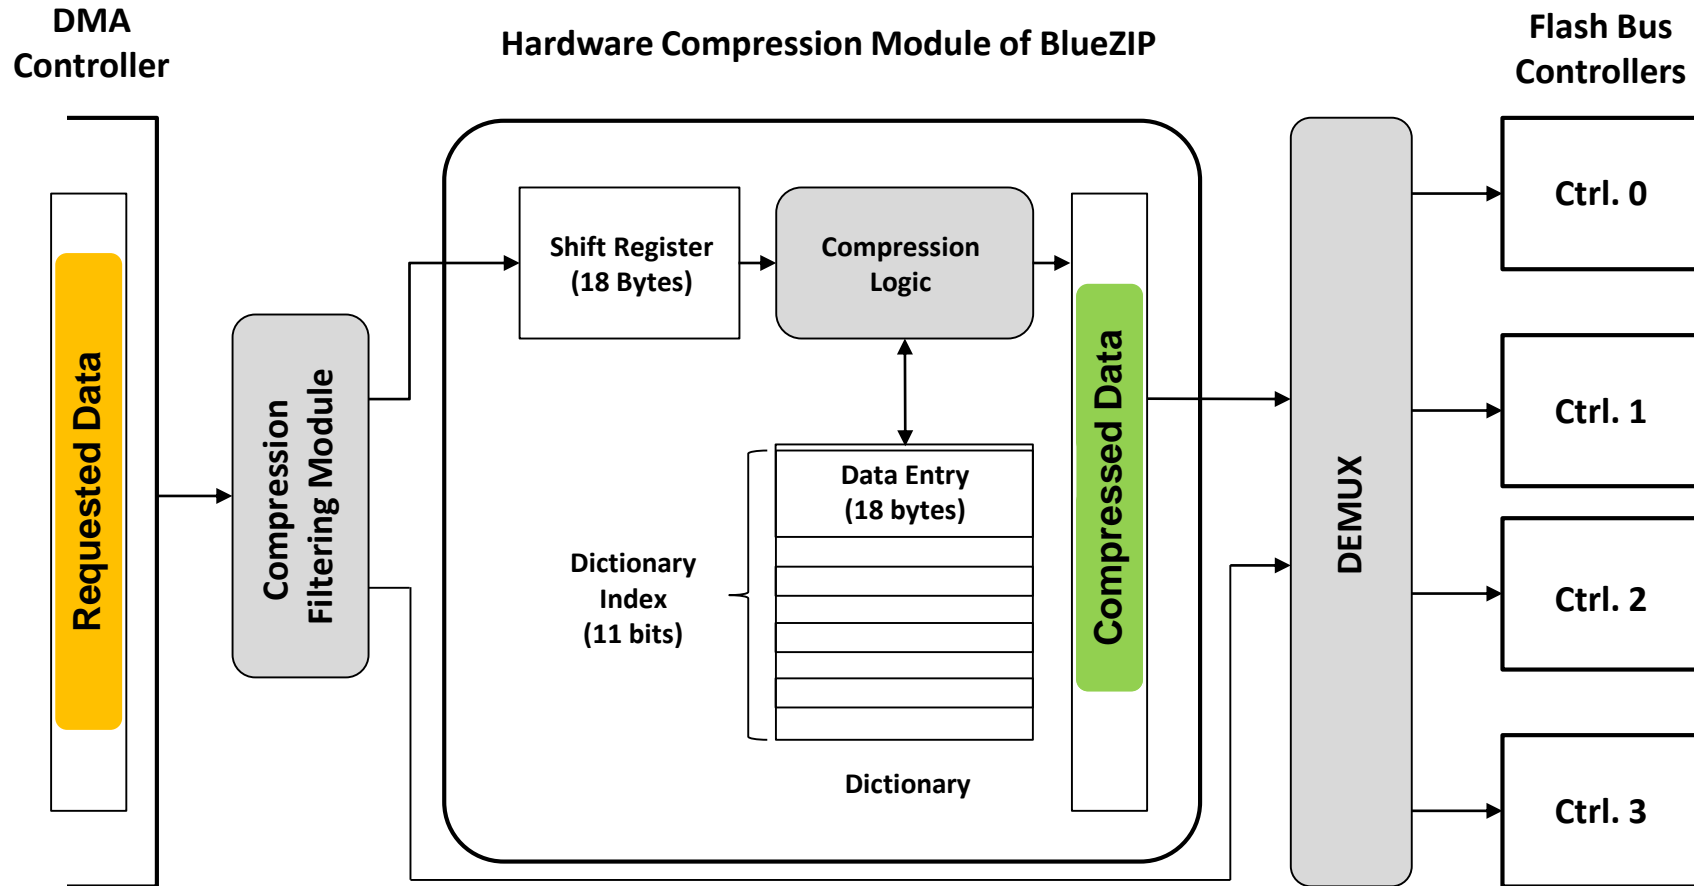
HW: Compression Algorithm

■ Use the LZRW3 compression algorithm

- Relatively high compression ratio
- Easy hardware implementation

Compression Algorithm	Hardware Complexity	Performance (Cycles)	Compression Ratio	Description
X-Match	Low	1,024 cycles / 4 KB (20.48 us@ 50 MHz)	Low (e.g., 20%)	Hardware-Based Memory Compression
LZ77	High	4,096 cycles > / 4 KB	High (e.g., 50%)	File Compression
LZRW3	Middle	4,096 cycles / 4 KB (81.92 us@ 50 MHz)	High (e.g., 40%)	File Compression

HW: Implementation



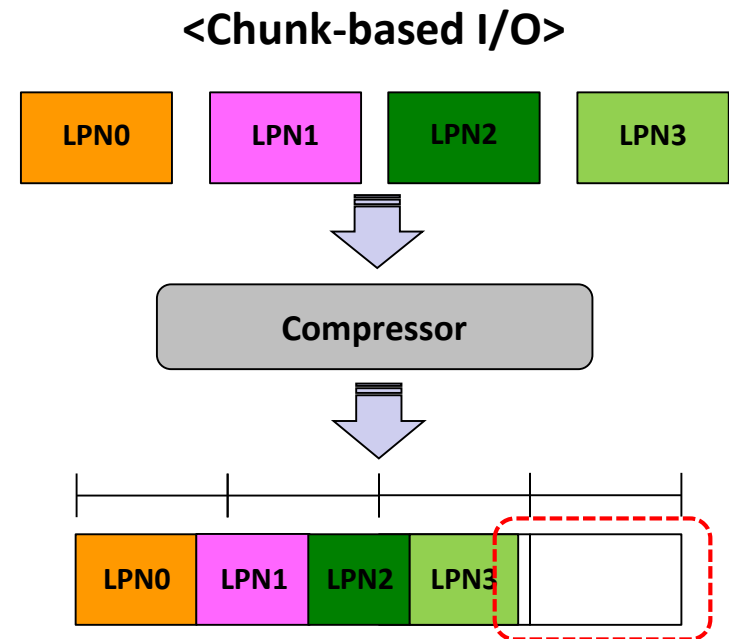

HW-SW: Compression Unit

■ Chunk-based I/O

- Compress several logical pages into several physical pages
- Mitigate the internal fragmentation problem

■ Read performance penalty

- Use a read buffer that holds previously decompressed pages



SW: Compression-Aware FTL

Write Buffer

301	302	507	508	...
-----	-----	-----	-----	-----

Page Mapping Table

Logical Page Address	Physical Page Address
⋮	⋮
301	2311
302	2311
507	2311
508	2311

Physical Page Address

⋮
2311
2312
2313

Data Chunk Table

Valid Page Counter	No. of Physical Pages	Compression Indicator
⋮	⋮	⋮
4	3	1
4	3	1
4	3	1
⋮	⋮	⋮

Physical Page Address:

Compression Chunk (3 pages)

...			2311	2312	2313			...
			301	302	507	508		

NAND Flash Memory


SW: Selective Compression

■ The data size expansion problem

- The size of compressed data > the size of original data
- Degrade storage performance and lifetime
- Usually observed in writing multimedia files

■ Selective compression

- Detect poorly compressed sequential writes in advance and do not compress those data



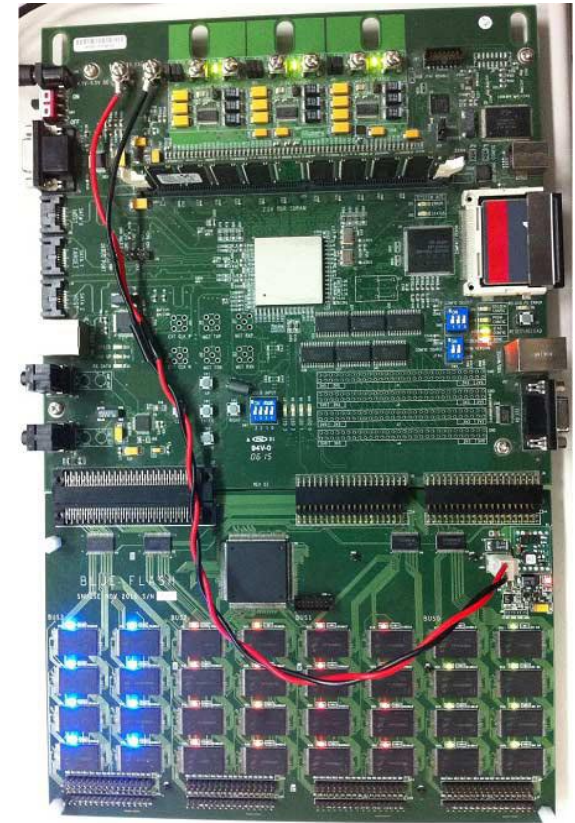
Experimental Settings

■ Benchmarks

Benchmark	Description	Compression Ratio
SENSOR	A set of sensor data files which were collected during a semiconductor fabrication process	Very high
LINUX	A subset of the Linux kernel 2.6.32 source files	high
DOCUMENT	A set of documents and image files	medium
MP3	A set of MP3 files already highly compressed	Very low

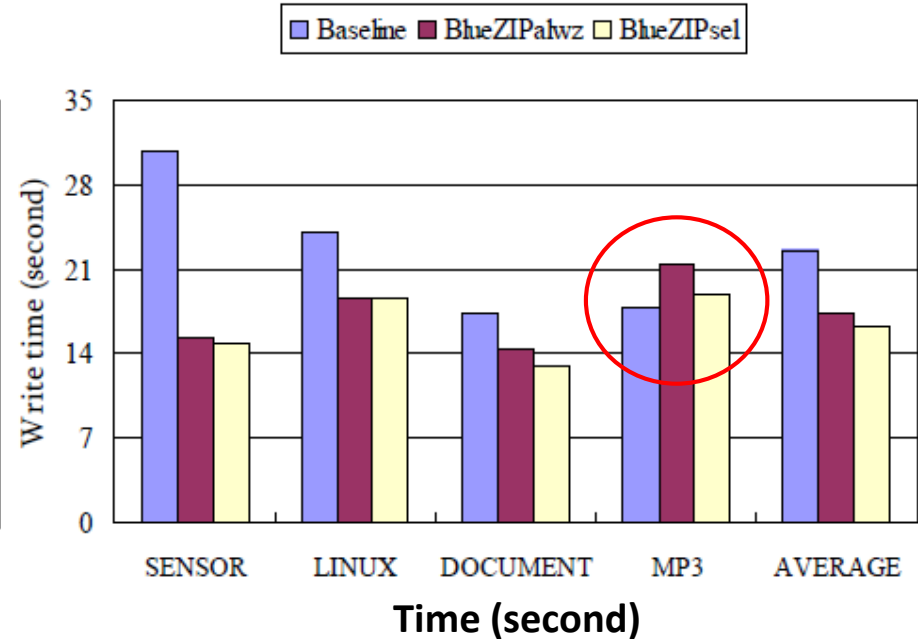
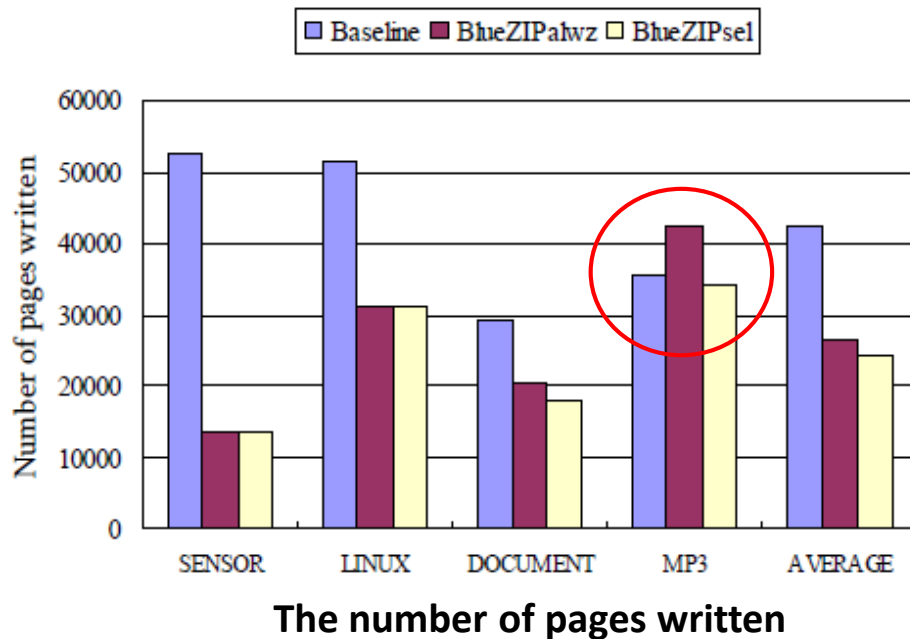
■ SSD Configurations

Configurations	Description
Baseline	Use no lossless compression
BlueZIP ^{alwz}	Use lossless compression all the time
BlueZIP ^{sel}	Use selective compression



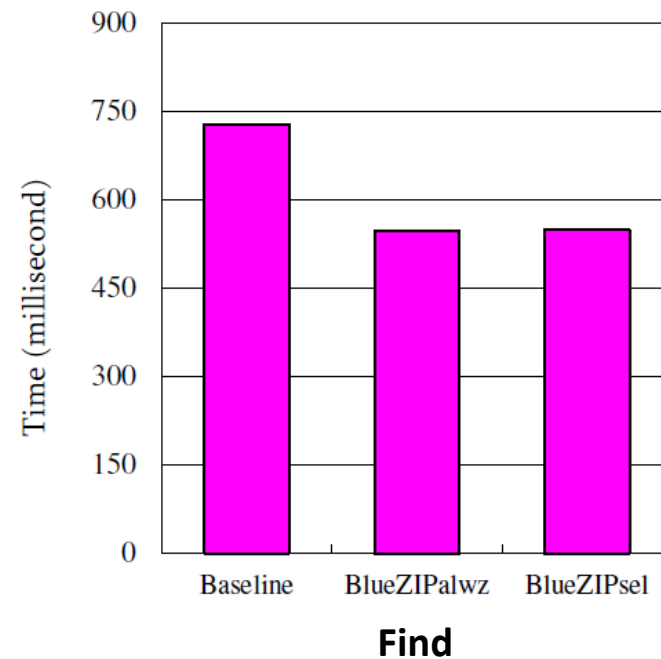
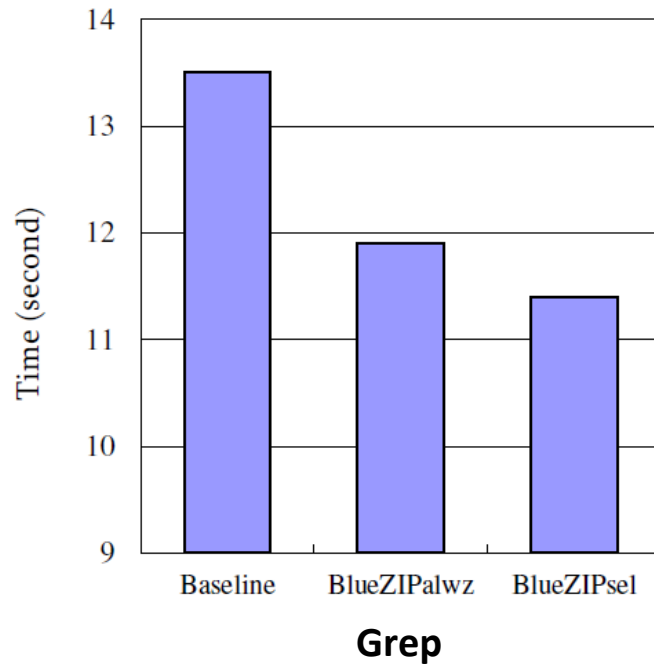
<BlueZIP prototype>

Experimental Results



- **BlueZIP^{sel} writes 38% less data over Baseline.**
 - The amount of written data is increased with BlueZIP^{alwz} due to the data expansion problem of lossless compression.
- **BlueZIP^{sel} achieves 17%-50% higher performance than Baseline**

Read Performance



- BlueZIPsel improves the overall read performance by 20% on average.
- The reduction in the number of pages sufficiently offsets the decompression overhead.

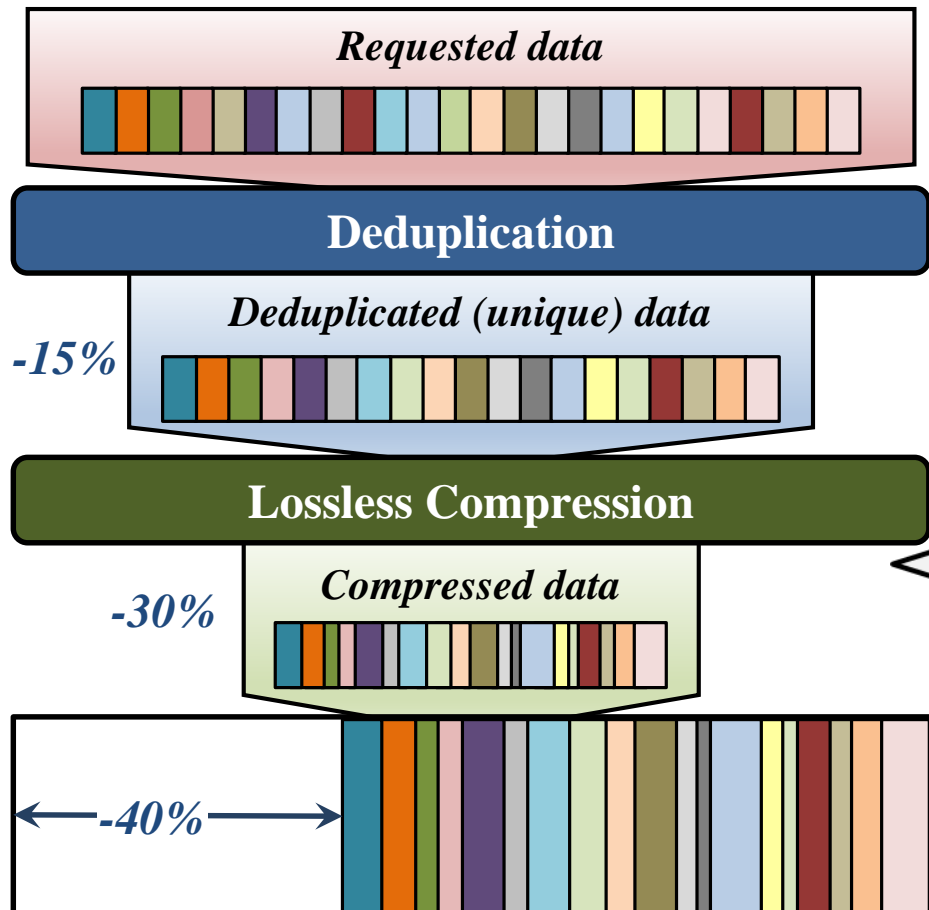
Outline

- Data Deduplication
- Data Compression
- **Case Studies**
 - CA-SSD: Content-aware Solid-state Drives
 - BlueZIP: Hardware-accelerated Compression
 - **DAC: Dedup-assisted Compression**

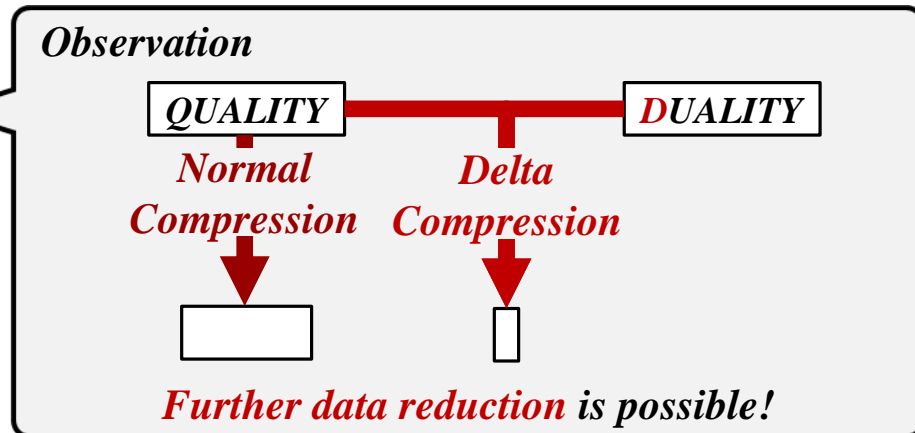
Integration of Capacity-optimization Techniques

■ Maximizing the data compression ratio

- Improvements can be accumulated

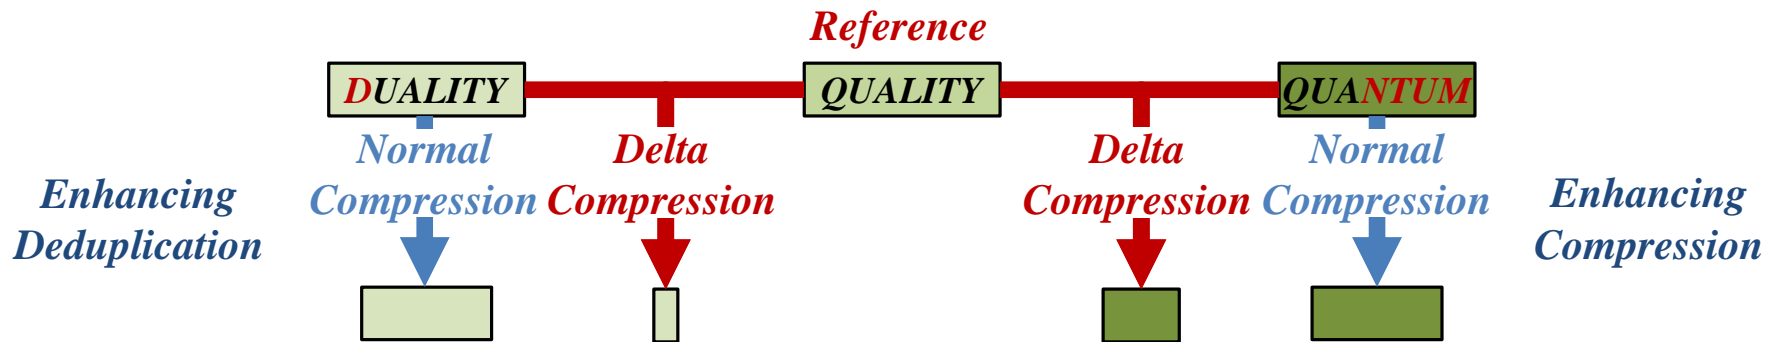


Q: Can we do better?




DAC: Dedup-Assisted Compression

- A novel integrated data reduction technique
 - Exploiting data similarity to further increase data reduction ratio



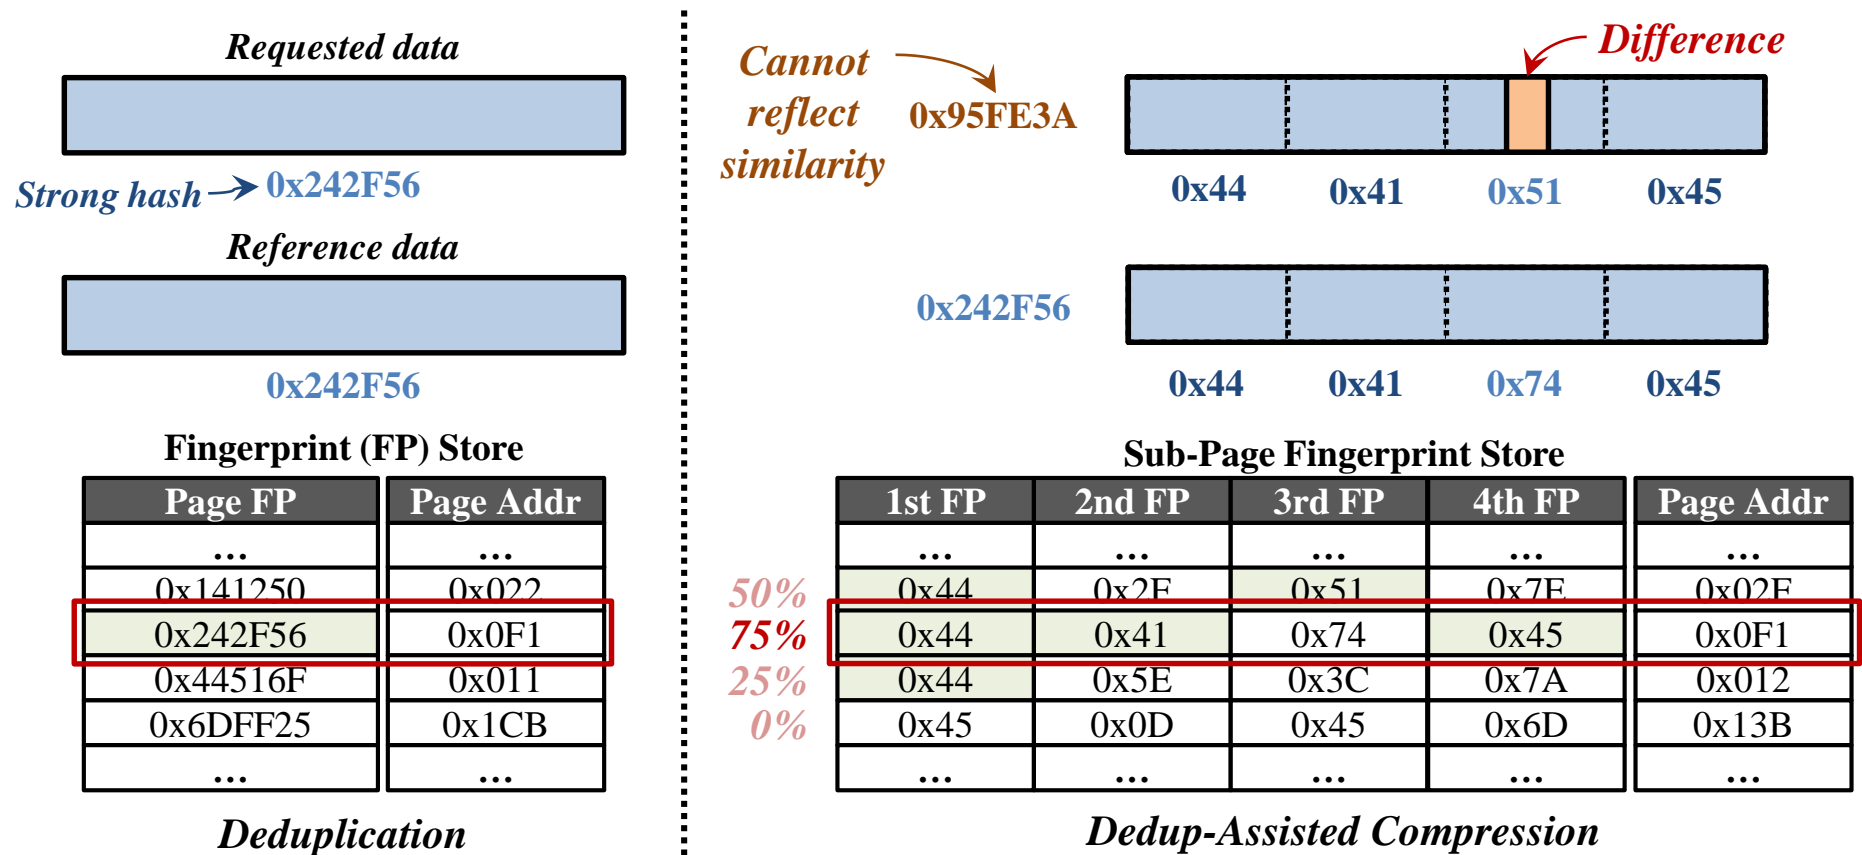
- Reference search is not easy



DAC: Reference Data Search

■ Exploiting partial fingerprints (hashes) of data

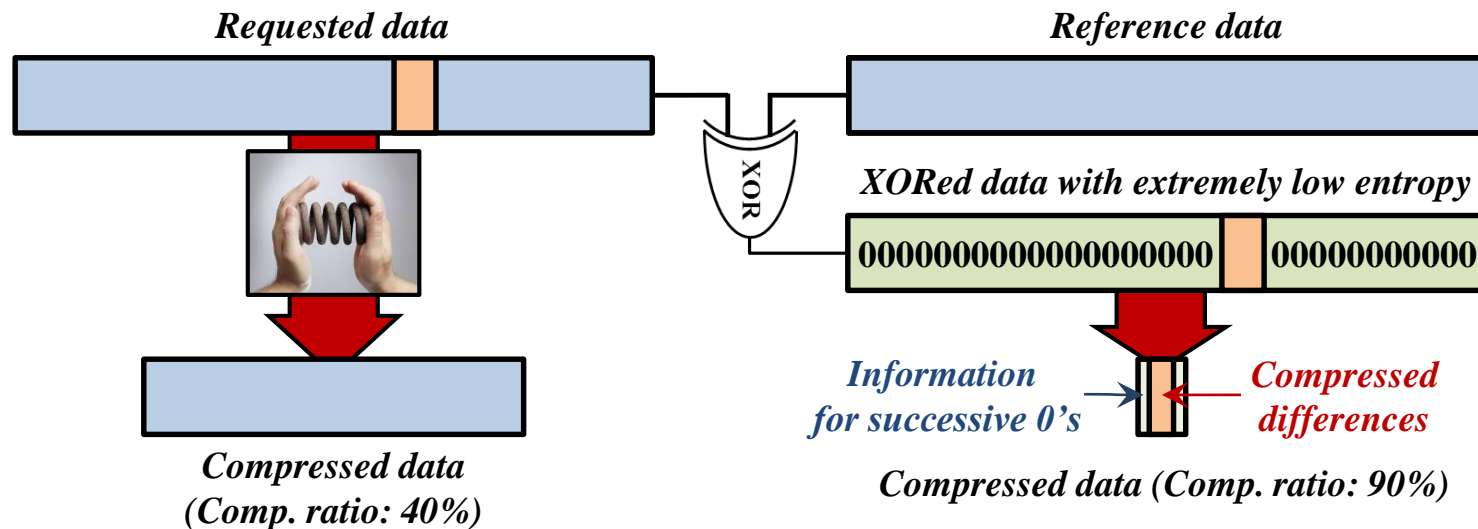
- Generated by the strong hash module for deduplication (Dedup-Assisted)



Wide searching range, constant searching time, and quantified data similarity

DAC: Data Compression with a Reference

■ Exploiting XOR logic to extremely reduce data entropy



■ Lossless compression algorithm: XmatchPRO

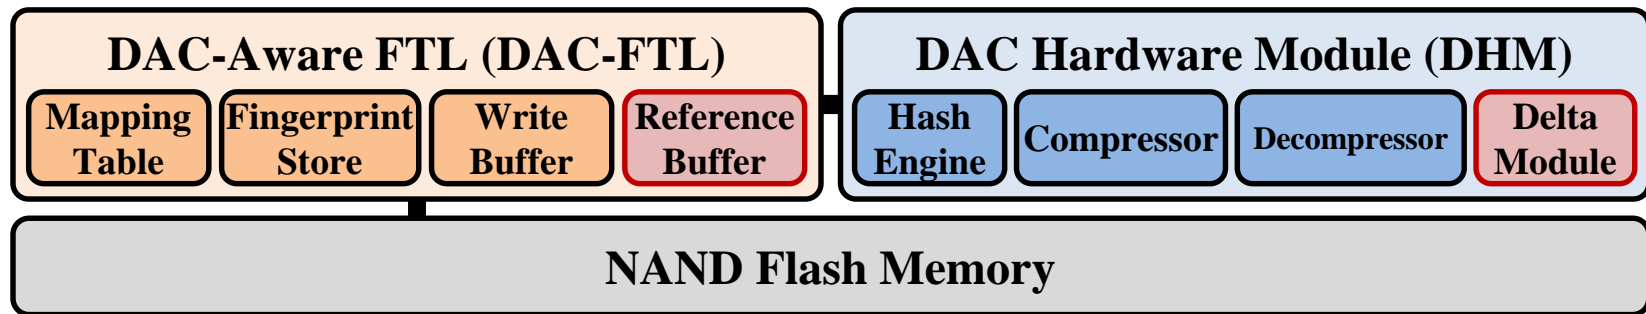
- Employing both the run-length and dictionary-based data encoding
- Easily implementable in H/W
- High performance and comparable compression ratio

Higher compression efficiency w/o an additional delta-aware data encoding

Experimental Environment

■ Target SSD Architecture

- Recent high-end SSDs employing deduplication and compression



■ Evaluation on a storage emulation environment

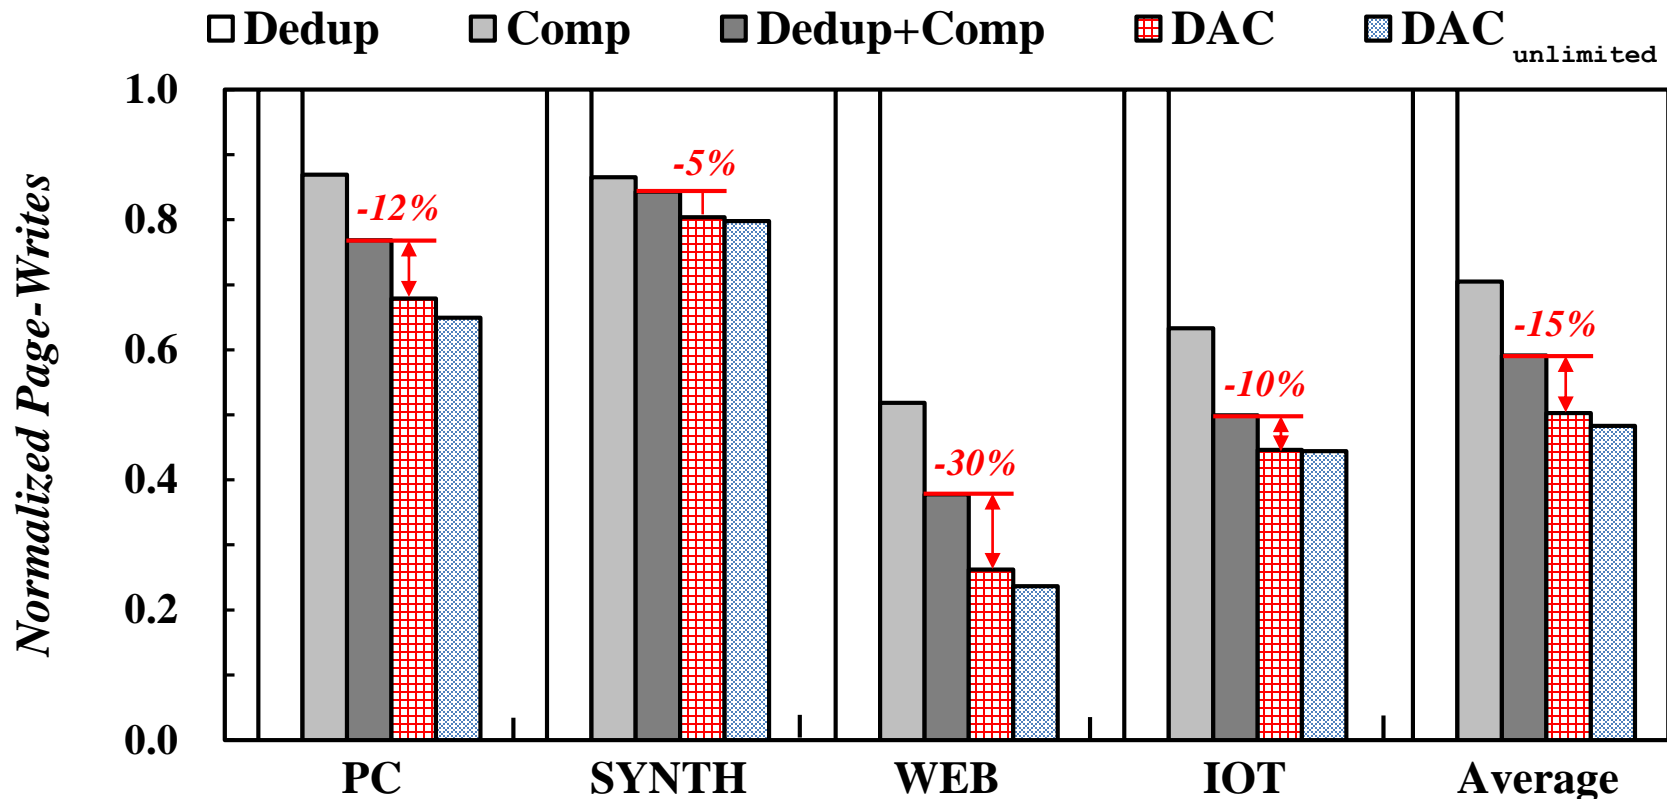
- Hardware modules for a strong hash function (MD5) and lossless compression (XmatchPRO) were emulated by software.

■ Benchmarks: block I/O traces collected from a high-end PC

Benchmark	PC	SYNTH	WEB	IOT
Data Redundancy	<i>Moderate</i>	<i>Low</i>	<i>Very High</i>	<i>High</i>
Data Compressibility	<i>Moderate</i>	<i>Low</i>	<i>Very High</i>	<i>High</i>

Evaluation Result

- Compared with a combination of deduplication and lossless compression, DAC reduces write traffic by up to 30% and 15% on average
- DAC can further increase data reduction whether data entropy is high or not



End of Chapter 10