

12. Persistent Memory

Special Topics in Computer Systems:
Modern Storage Systems
(IC820-01)

Instructor:

Prof. Sungjin Lee (sungjin.lee@dgist.ac.kr)

Outline

- **What is Persistent Memory**
- **Memory Mode**
- **App Direct Mode**

Persistent Memory

■ Byte addressable stable memory devices

■ *Phase Change (PCM)*

- Crystalline material heated to two different phases

■ *Spin Torque Transfer (STT-RAM, MRAM)*

- Magnetic spin captured in ferrous element

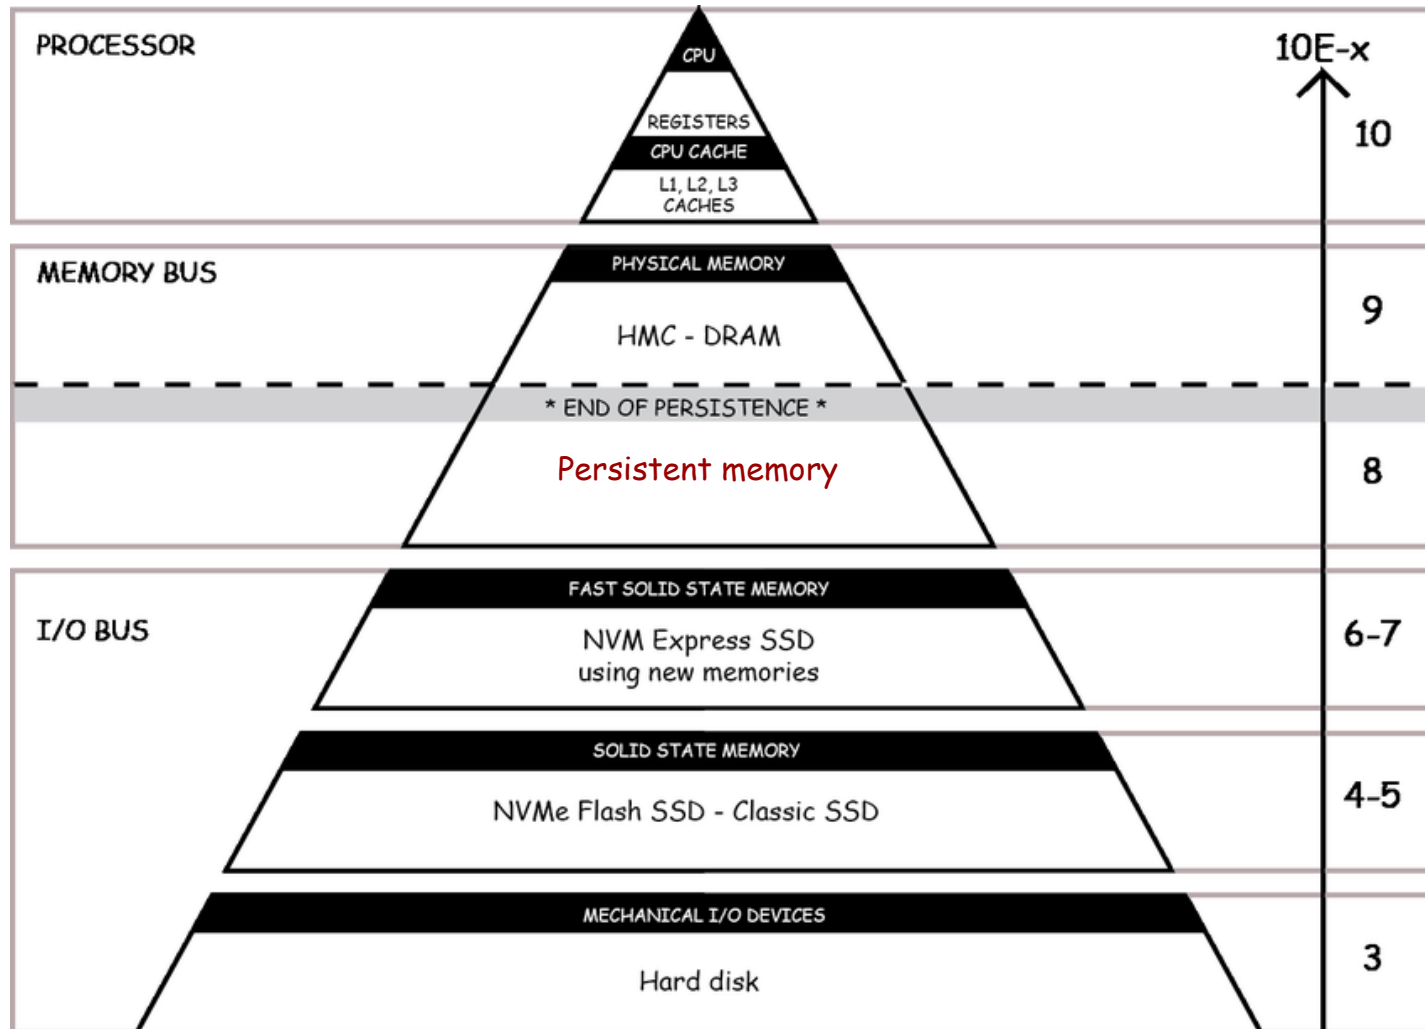
■ *ReRAM*

- Resistive RAM

■ Persistent memory will fit a niche between DRAM and NAND

- Power advantages over DRAM because there is no “refresh” cycle necessary to preserve logic value
- Cost disadvantages over NAND because block-addressing makes NAND more compact

System Hierarchy with Persistent Memory



Intel's Optane (or 3D X-Point)

■ New Technology from Intel/Micron

- Similar to Phase Change and Resistive Memories
- Bit addressable
- No Erase requirement

■ '~1000 faster' than NAND

- But still 3x to 5x slower than DRAM


■ Process has a trajectory to catch NAND density

- This is a big deal, first parts are **128, 256, and 512 GB!**

■ Much improved durability

- However, even at 100 million write cycles, you could wear it out in seconds w/out wear leveling

■ Products are memory (DIMM) or storage devices (PCIe)



Intel's Optane (or 3D X-Point) (Cont.)



■ Optane™ DC Persistent Memory Module (Optane DC PMM)

Memory attributes

Performance comparable to DRAM at low latencies

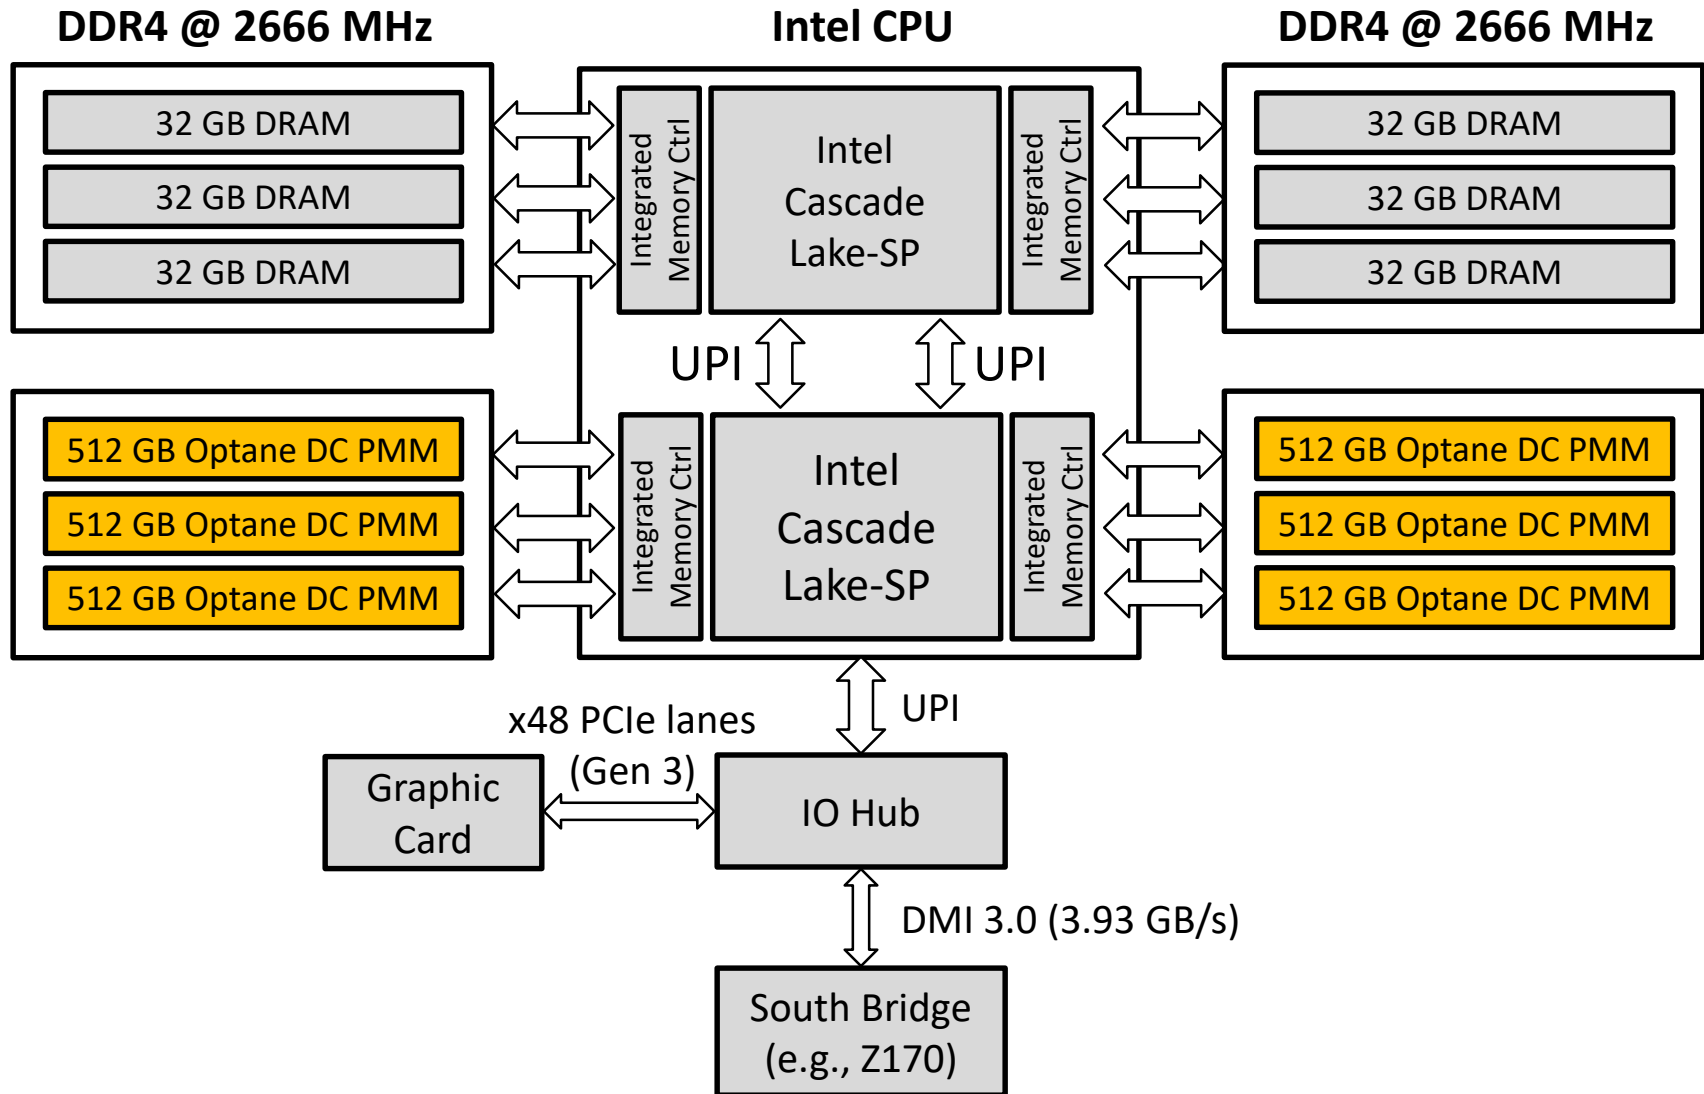
Storage attributes

Data persistence with Higher capacity than DRAM




	DRAM	Optane DC PMM (6x)
Read Latency	81 ns	306 ns
Write Latency	86 ns	94 ns
Read Throughput	120 GB/s	39.4 GB/s
Write Throughput	80 GB/s	13.9 GB/s

Storage Bandwidth Hierarchy



Intel Xeon Scalable with 6TB of Optane DC

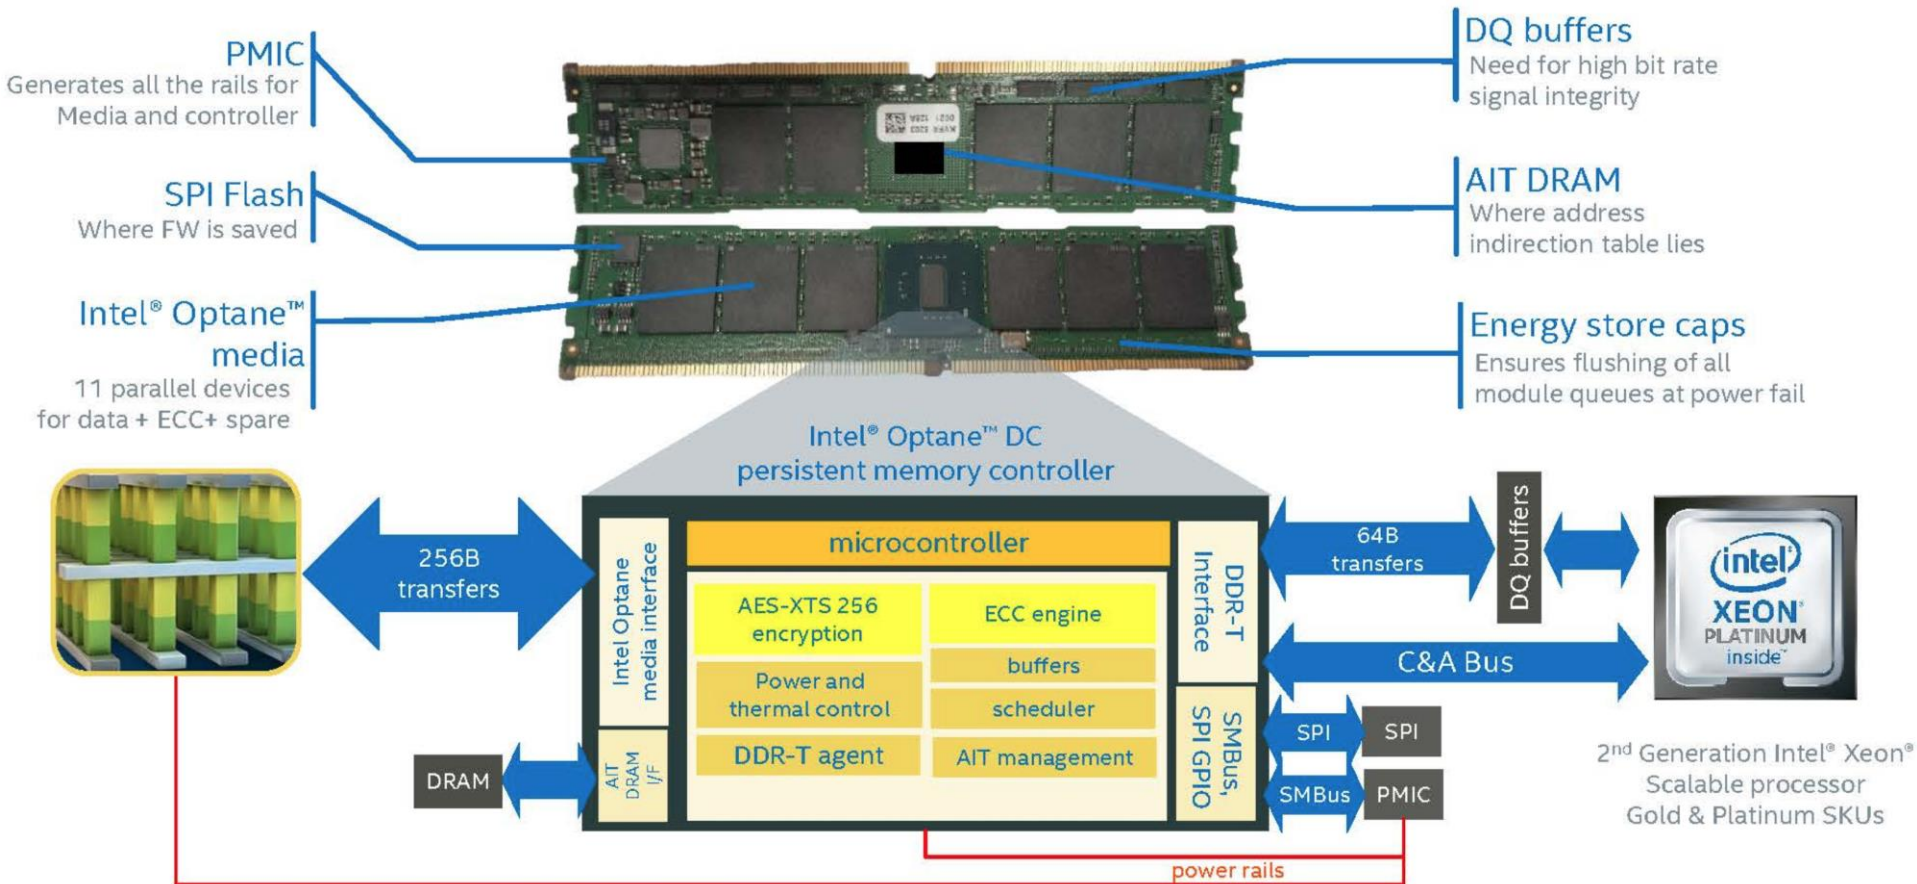



```
Handle 0x009C, DMI type 17, 40 bytes
Memory Device
  Array Handle: 0x0092
  Error Information Handle: Not Provided
  Total Width: 72 bits
  Data Width: 64 bits
  Size: 32 GB
  Form Factor: DIMM
  Set: None
  Locator: CPU4_DIMM_F1
  Bank Locator: NODE 8
  Type: DDR4
  Type Detail: Synchronous
  Speed: 2666 MT/s
  Manufacturer: Samsung
  Serial Number: 404B0F29
  Asset Tag: CPU4_DIMM_F1_AssetTag
  Part Number: M393A4K40CB2-CTD
  Rank: 2
  Configured Clock Speed: 2666 MT/s
  Minimum Voltage: 1.2 V
  Maximum Voltage: 1.2 V
  Configured Voltage: 1.2 V
```

```
Handle 0x009E, DMI type 17, 40 bytes
Memory Device
  Array Handle: 0x0092
  Error Information Handle: Not Provided
  Total Width: 72 bits
  Data Width: 64 bits
  Size: 258496 MB
  Form Factor: DIMM
  Set: None
  Locator: CPU4_DIMM_F2
  Bank Locator: NODE 8
  Type: <OUT OF SPEC>
  Type Detail: Synchronous Non-Volatile
  Speed: 2666 MT/s
  Manufacturer: Intel
  Serial Number: 00001817
  Asset Tag: CPU4_DIMM_F2_AssetTag
  Part Number: NMA1XBD256GQS
  Rank: 1
  Configured Clock Speed: 2666 MT/s
  Minimum Voltage: 1.2 V
  Maximum Voltage: 1.2 V
  Configured Voltage: 1.2 V
```


CPU recognizes it
as **256 GB DIMM DRAM**
(Memory mode)

Optane Architecture



DDR-T Protocol

- **DDR protocol on top of electrical/mechanical interface for DDR4**
 - DDR DRAM operate synchronously which are not suitable for Optane media
- **DDR-T allow for *asynchronous* cmd/data timing**
 - Controller uses request/grant scheme to communicate with host controller
 - Data bus direction and timing controlled by host
 - Command packet per request sent from host to Optane DC PMM
- **Transaction can be re-ordered**
- **64B cache line access granularity (similar to DDR4)**



Address Indirection Table (AIT)

- **Optane DC PMM media has limited write endurance like NAND flash**
 - Internal address translation is necessary for wear-leveling and bad-block management

- ***Address Indirection Table (AIT)***
 - Translate the DIMM physical address to an internal Optane DC media device address (similar to what FTLs do in SSDs)
 - The AIT table resides in Optane media, though on-DIMM DRAM keeps a copy of the AIT entries
 - The granularity of Optane DC media is 256 bytes, while a cache line size is 64 bytes
 - Write amplification occurs – smaller stores are handled as read-modify-write operations by the controller

Outline

- **What is Persistent Memory**
- **Memory Mode**
- **App Direct Mode**

Two Modes of Optane DC PMM

■ 1. *Memory Mode*


- DDR4 DIMMs operate as caches for slower Optane DC PMM
 - Direct mapped write-back cache with 64B cache lines
- Optane DC PMMs are exposed as a large “*volatile*” memory
- Do not have persistence

■ 2. *App Direct Mode*

- Optane DC PMMs are directly exposed to the CPU and OS as “persistent” storage
- Both Optane DC PMMs and DDR4 DIMMs are visible to OS as memory devices
- Optane DC PMM is exposed as configurable regions on contiguously-addressed ranges

Memory Mode Details

- No software/application changes required
- To mimic traditional memory, data is “volatile”
- DRAM is “*near memory*”
 - Used as a write-back cache
 - Managed by host memory controller
 - Within the same host memory controller, not across
- Optane DC PMM is “*far memory*”
 - Managed by host software
- Ratio of far/near memory (PMEM/DRAM) can vary



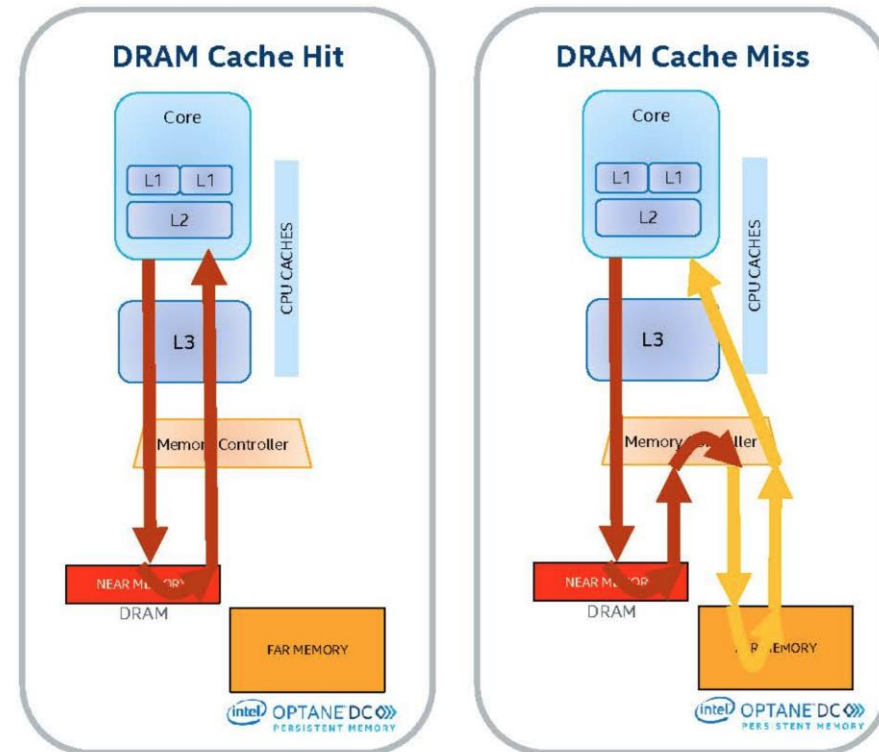
Memory Mode Details (Cont.)

■ Good locality means near-DRAM performance

- Cache hit: latency same as DRAM
- Cache miss: latency DRAM + Optane DC persistent memory


■ Performance varies by workload

- Best workloads have the following traits:
 - Good locality for high DRAM cache hit rate
 - Low memory bandwidth demand
- Other factors:
 - # of reads > # of writes



Memory Mode Performance

- Order of magnitude lower latency than SSD
- 2x read/write bandwidth vs disk, with one module, more with multiple modules




Memory Mode Performance (Cont.)

- **Optane DC PMM is programmable for different power limits for power/performance optimization**
 - 12W – 18W, in 0.25 watt granularity – for example: 12.25W, 14.75W, 18W
 - Higher power settings give best performance
- **Performance varies based on traffic pattern**
 - Contiguous 4 cache lines (256B) vs. single random cache line (64B)
 - Reads vs. writes

Granularity	Traffic	Module	Bandwidth
256B (4x64B)	Read	256GB, 18W	8.3 GB/s
256B (4x64B)	Write		3.0 GB/s
256B (4x64B)	2 Read/1 Write		5.4 GB/s
64B	Read		2.13 GB/s
64B	Write		0.73 GB/s
64B	2 Read/1 Write		1.35 GB/s

Memory Mode Performance (Cont.)

■ Performance with SPEC CPU 2006 and 2017



Outline

- **What is Persistent Memory**
- **Memory Mode**
- **App Direct Mode**

Two Modes of Optane DC PMM

■ 1. *Memory Mode*

- DDR4 DIMMs operate as caches for slower Optane DC PMM
 - Direct mapped write-back cache with 64B cache lines
- Optane DC PMMs are exposed as a large “*volatile*” memory
- Do not have persistence

■ 2. *App Direct Mode*

- Optane DC PMMs are directly exposed to the CPU and OS as “persistent” storage
- Both Optane DC PMMs and DDR4 DIMMs are visible to OS as memory devices
- Optane DC PMM is exposed as configurable regions on contiguously-addressed ranges

App Direct Mode Details

■ PMEM-aware software/application required

- Adds a new tier between DRAM and block storage
- Industry open standard programming model (e.g., PMDK)

■ In-place persistence

- No paging, context switching, interrupts, nor kernel code executes


■ Byte addressable like memory

- Load/store access, no page caching

■ Cache coherent

■ Ability to do DMA & RDMA

- No copy to DRAM for data transfer to block storage



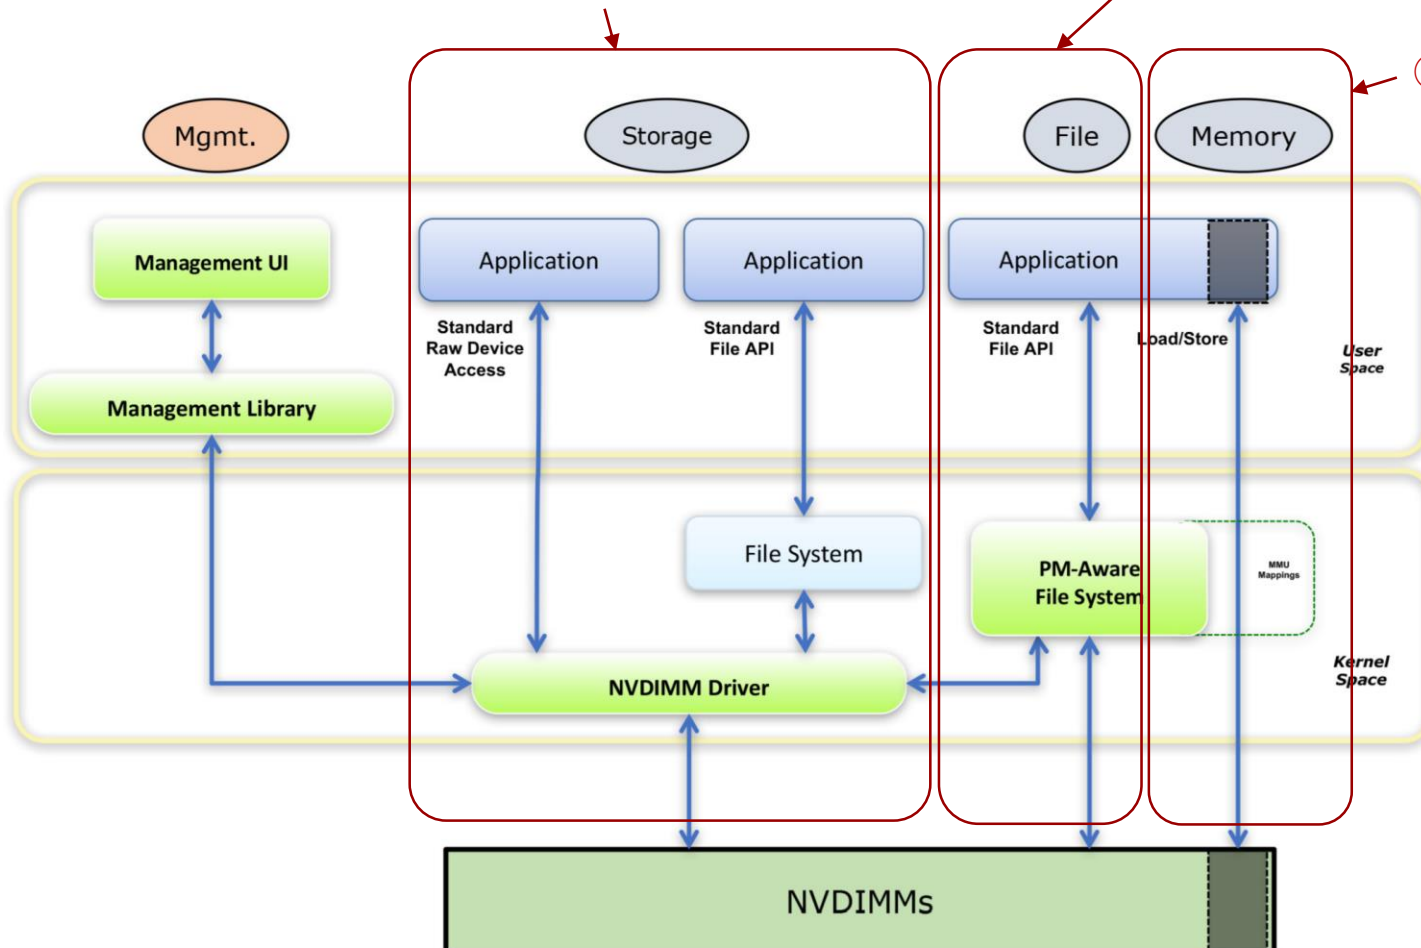
App Direct Mode Details (Cont.)

■ Persistent memory programming model

② Persistent memory-aware file system

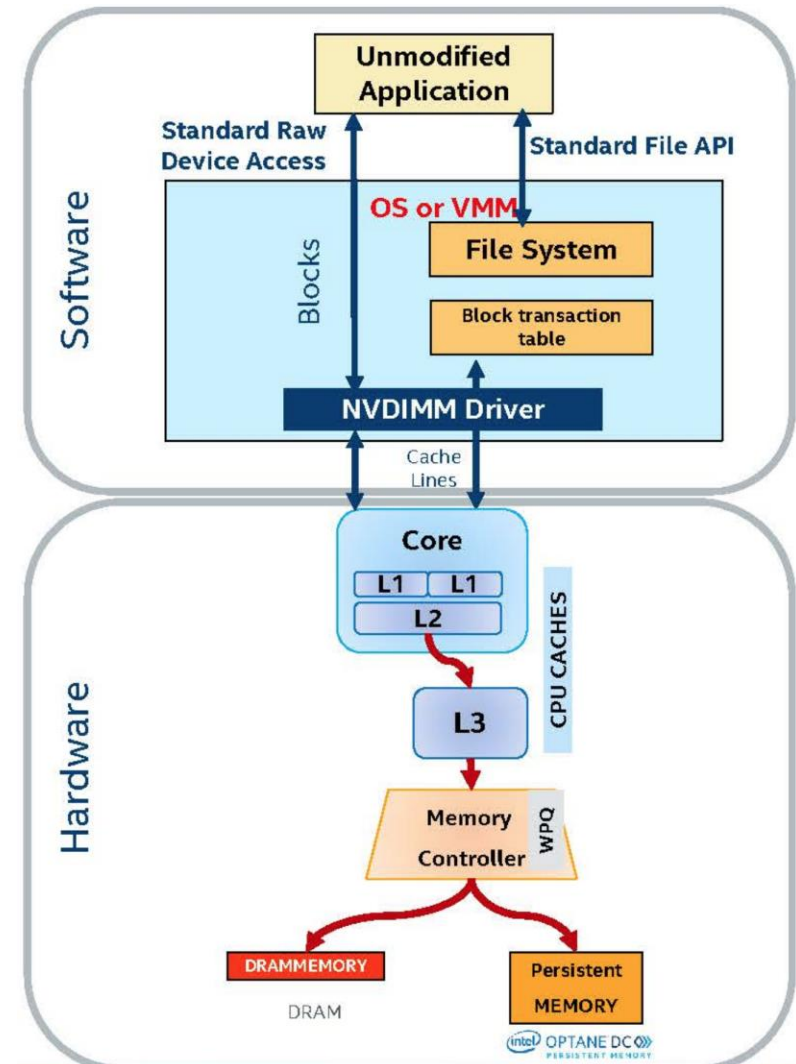
① Traditional block storage

③ load & store access



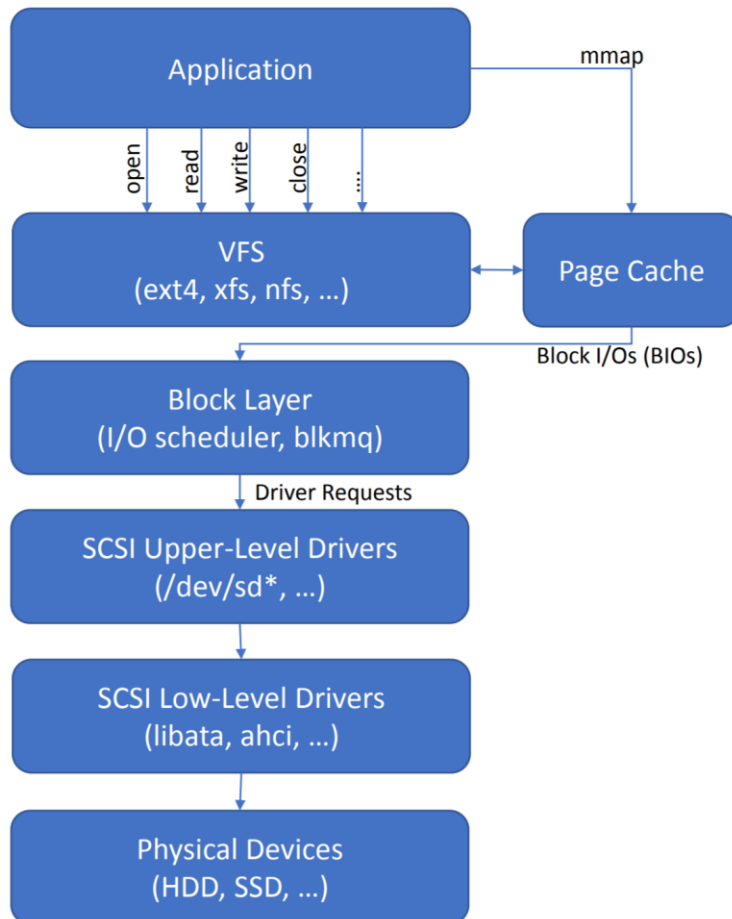
Traditional Block Storage

- **Operates in blocks like SSD/HDD**
 - Traditional read/write instructions
 - Work with existing file systems
- **NVDIMM driver required**
 - Support starting Linux kernel 4.2
- **Scalable capacity**
- **Higher endurance than SSDs**
- **High performance block storage**
 - Low latency, higher bandwidth, high IOPS

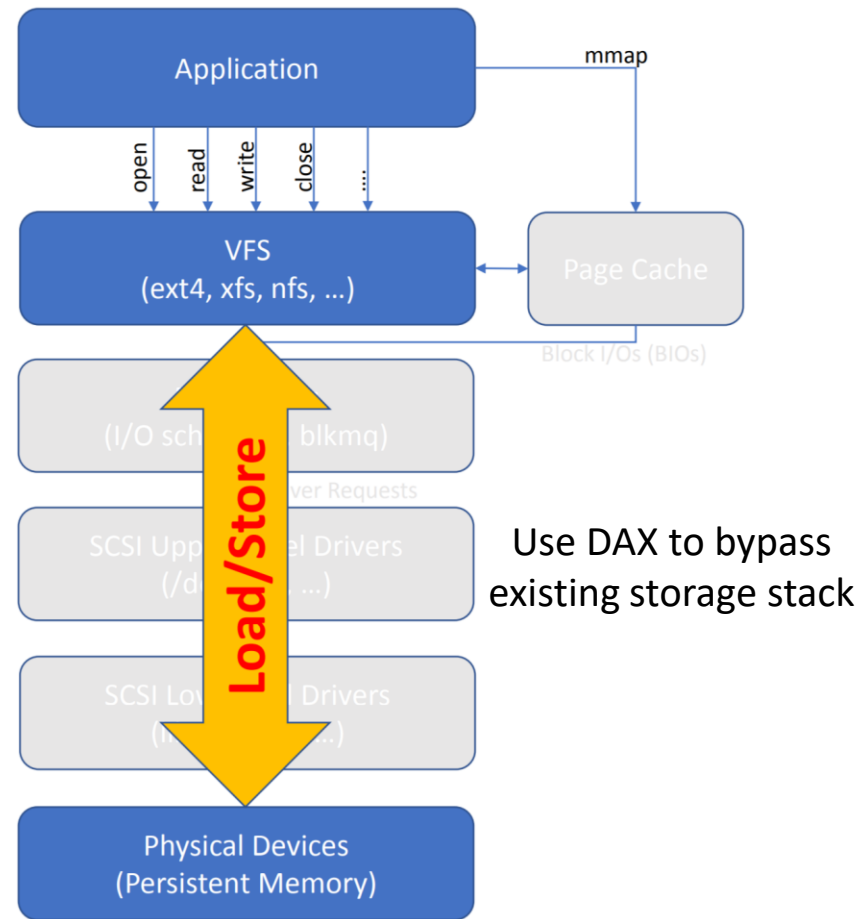


Persistent Memory-aware File System

■ Traditional Storage Stack



■ Persistent Memory-aware Stack




DAX-enabled File System

■ What is DAX (*Direct Access*)

- Allow file systems to directly access to persistent memory without using the system page cache
- Added to Linux and Windows operating systems

■ Why is DAX necessary?

- Page Cache
 - Mediate between fast memory and slow storage
- Elegant design
 - Read-ahead
 - Write-back policy
 - ...
- Valid for decades until now,
 - Becoming a new bottleneck!



DAX-enabled File System (Cont.)

■ Use existing `mmap` semantics

- The persistent memory is directly mapped to the user space

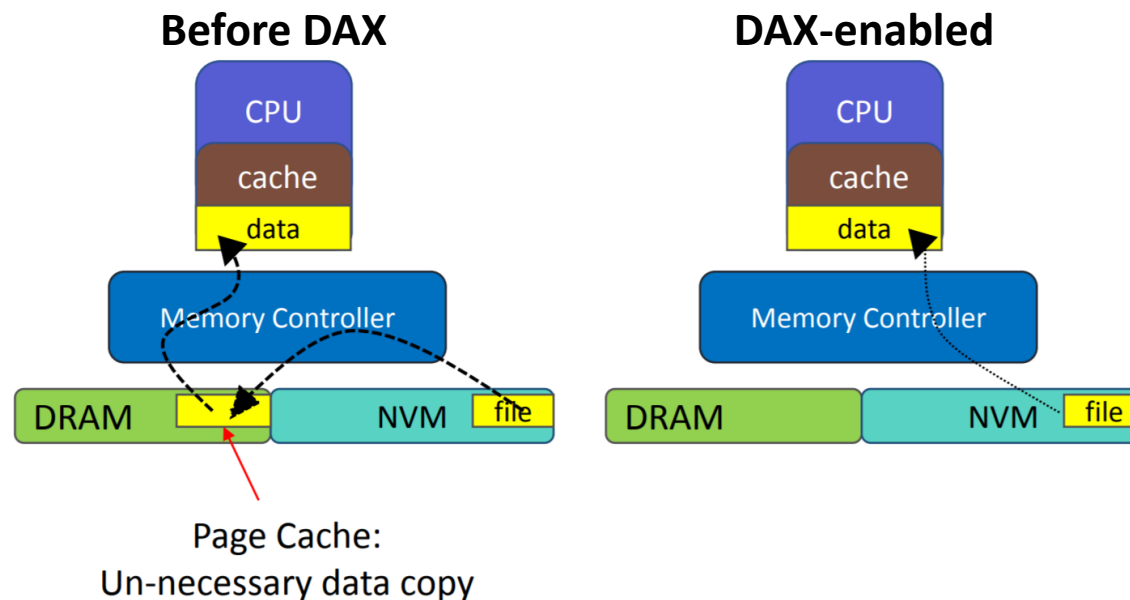
■ In-place update

- All the data is directly written to the persistent memory without buffering in the page cache

■ True device performance

■ File system implementation with DAX

- EXT4-DAX and NTFS-DAX



Is DAX all enough?




- **Simple answer: No**
- **More details?**
 - Crash consistency not guaranteed!

Crash Consistency

■ Definition

- “Ensure that the file system keeps the on-disk image in a reasonable state given that crashes can occur at arbitrary points in time.”

 Crash!
`strcpy(pmem, "Hello World!");`

 Hello
 Hello World!
 *Nothing*

How about using `fsync()` or `msync()`

- Some might want to `fsync()` or `msync()` to ensure the changes are persistence

```
strcpy(pmem, "Hello World!");  
msync(pmem, 12, MS_SYNC);
```

- Do they work with the persistent memory?

- Answer: No
- Why: `fsync()` and `msync()` are used to flush out dirty pages in the page cache to storage devices. But, with DAX, no page cache is used

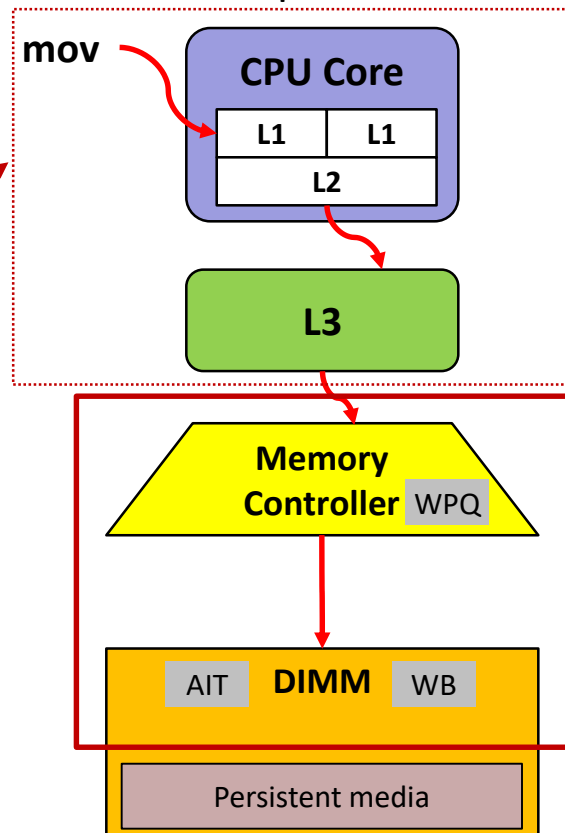
Persistence Domain

■ Asynchronous DRAM refresh (ADR) (Intel's feature)

- Ensure the data is in a “safe” state during a power loss event or system crash
- Triggers a hardware interrupt to the memory controller
- The controller will *flush* the data buffers and place the DRAM in *self-refresh*
- *It does not flush the CPU cache*

X86 caches are quite large;
Take more energy than the capacitors
in a power supply can provide

↓
Software makes sure that
data is flushed to persistent domain
using
CLFLUSH or CLWB



Persistence Domain:
Minimum required power
fail protection domain

X86 Cache Flush Instructions

- **Simply executing a store instruction is not enough to make data persistent**
 - The data may be sitting in the CPU caches and could be lost by a power failure
- **Additional cache flush instructions (e.g., CLFLUSHOPT and CLWB) are required to make the stores persistent**

CLFLUSH	This instruction, supported in many generations of CPU, flushes a single cache line. Historically, this instruction is serialized, causing multiple CLFLUSH instructions to execute one after the other, without any concurrency.
CLFLUSHOPT	This instruction, newly introduced for persistent memory support, is like CLFLUSH but without the serialization
CLWB	Another newly introduced instruction, CLWB stands for cache line write back. The effect is the same as CLFLUSHOPT except that the cache line may remain valid in the cache, but no longer dirty
SFENCE	Ensure the flushes are complete before continuing

CFLUSH

■ Order writes by flushing cachelines via CLFLUSH

STORE data[0] = 0xFOOD

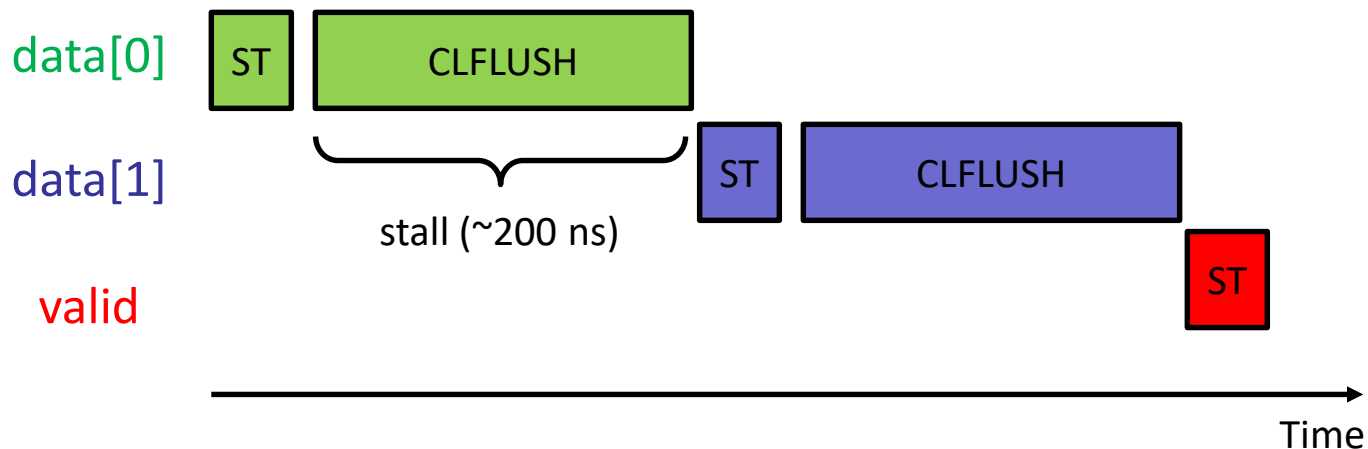
STORE data[1] = 0xBEEF

CLFLUSH data[0]

CLFLUSH data[1]

STORE valid = 1 /* mark that data[0] and data[1] are valid */

■ CLFLUSH stalls the CPU pipeline and serializes execution



CLFLUSHOPT

- Provides unordered version of CLFLUSH
- Supports efficient cache flushing

STORE data[0] = 0xFOOD
STORE data[1] = 0xBEEF
CLFLUSHOPT data[0]
CLFLUSHOPT data[1]
SFENCE
STORE valid = 1

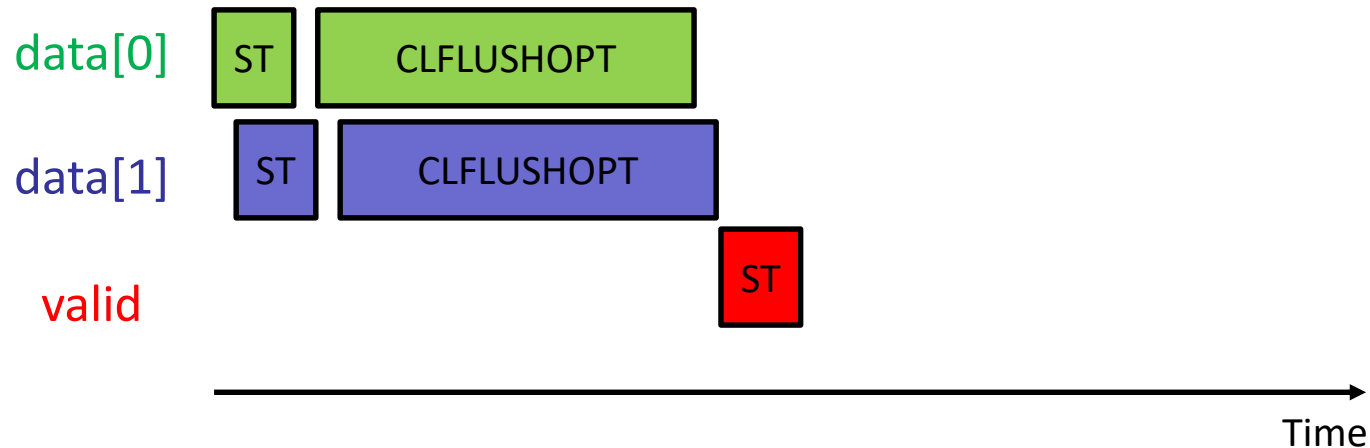
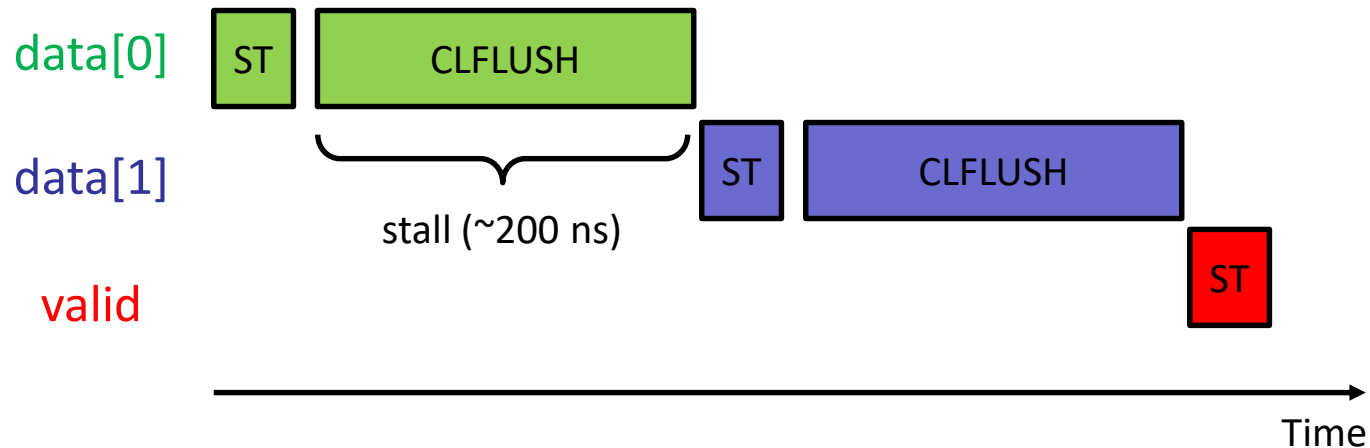
← Implicit orderings

← Explicit ordering point

The diagram illustrates the execution order of the code sequence. Two curved arrows, one green and one blue, originate from the first two lines of code and point to the third and fourth lines, respectively, indicating implicit orderings. A straight arrow points from the text 'Explicit ordering point' to the 'SFENCE' line, indicating that the fence enforces a total order between the preceding and succeeding instructions.

CLFLUSHOPT (Cont.)

- Provides unordered version of CLFLUSH
- Supports efficient cache flushing

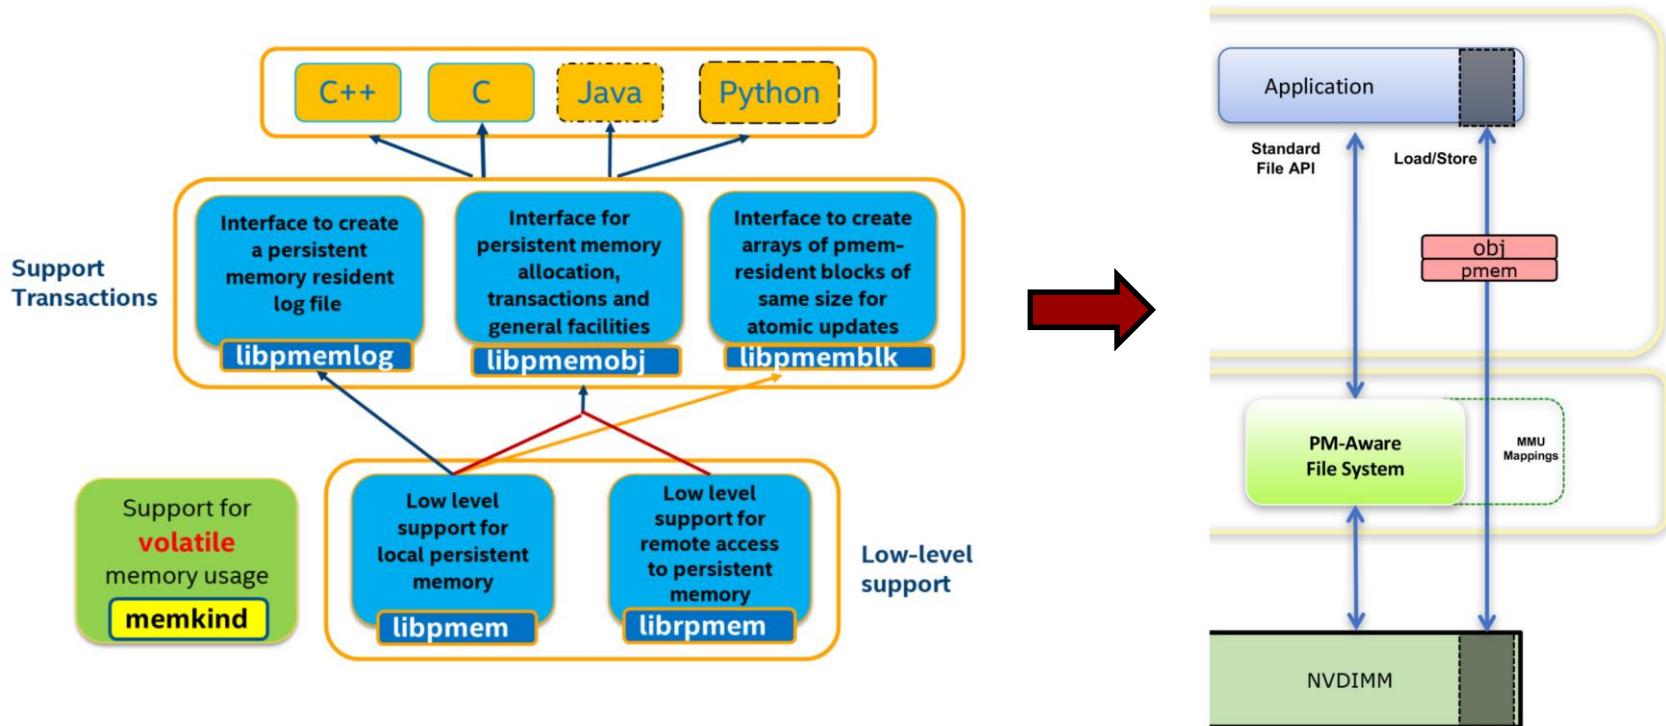


CLWB

- **Write backs modified data of a cacheline**
- **Does not invalidate the line from the cache**
 - Marks the line as non-modified
 - Improve cache hit ratios

The NVM Libraries

- A suite of six libraries for convenient use of persistent memory in user-space applications



libpmem: Basic Persistence Support

■ libpmem provides low level persistent memory supports

- Open and allocate the persistent memory
- Memory map the persistent memory to the user space
- Detect which types of flush instructions are supported by the CPU
- Performance-tuned routines for copying ranges of persistence memory

■ ...

```

57      /* create a pmem file */
58      if ((fd = open("/pmem-fs/myfile", O_CREAT|O_RDWR, 0666)) < 0) {
59          perror("open");
60          exit(1);
61      }
62
63      /* allocate the pmem */
64      if ((errno = posix_fallocate(fd, 0, PMEM_LEN)) != 0) {
65          perror("posix_fallocate");
66          exit(1);
67      }
68
69      /* memory map it */
70      if ((pmemaddr = pmem_map(fd)) == NULL) {
71          perror("pmem_map");
72          exit(1);
73      }
74      close(fd);

```

```

82      /* flush above strcpy to persistence */
83      if (is_pmem)
84          pmem_persist(pmemaddr, PMEM_LEN);
85      else
86          pmem_msync(pmemaddr, PMEM_LEN);

```

libpmemobj:

General-Purpose Allocations and Transactions

- libpmemobj allows persistent memory objects to be allocated in a way that is power fail safe
- Allows making an arbitrary number of changes atomic by encompassing the changes in a transaction
- Multithread safe and optimized for multithread scalability

```
class pmem_queue {
    /* entry in the list */
    struct pmem_entry {
        persistent_ptr<pmem_entry> next;
        p<uint64_t> value;
    };
    /* ... */
}
```

```
void
push(pool_base &pop, uint64_t value)
{
    transaction::exec_tx(pop, [&] {
        auto n = make_persistent<pmem_entry>();
        n->value = value;
        n->next = nullptr;
        if (head == nullptr) {
            head = tail = n;
        } else {
            tail->next = n;
            tail = n;
        }
    });
}
```

libpmemblk and libpmemlog:

Support for Specific Use Cases

■ **libpmemblk**

- Maintain a large array of persistent memory blocks, all the same size
- This is useful when an application is managing a block cache
 - The block size provided by the library is flexible, supporting blocks 512-bytes and larger

■ **libpmemlog**

- For a specific use case where the application frequently appends to a private log file, one that is read rarely, like during crash recovery

libmemkind:

The Volatile Use of Persistent Memory

■ Volatile use of persistent memory

- Places some data structures in persistent memory to avoid a large DRAM footprint, but doesn't really care that the memory is persistent

End of Chapter 12