

Project #2: DGIBox

Computer Algorithm
SE380

Instructor:

Prof. Sungjin Lee (sungjin.lee@dgist.ac.kr)


Project #2 Topic

- Develop an algorithms providing basic operations (write and read) that operate on a magical box, called a **DGIBox**
- Use all of the knowledge you have learned from this class
 - Sorting
 - Trees – Binary Search Tree, B-Tree, ...
 - Hash

What is DGIBox?

■ DGIBox is composed of a large number n of small buckets

- Each bucket looks like an array, indexed by an integer number
- You select any bucket in DGIBox and use it (e.g., inserting a key/value, ...)




DGIBox with n buckets, each of which can hold 4 items

- Each Bucket provides simple three operations:
 - **Set** (bid, index, value)
 - value=**Get** (bid, index)
 - **Empty** (bid)

How to Use DGIBox?

■ Example



DGIBox with n buckets, each of which can hold 4 items

```
Set(0, 0, 'A');
Get(0, 0); /* return 'A' */
Set(1, 2, 'C');
Get(0, 1); /* return nothing */
Set(1, 3, 'E');
Get(1, 2); /* return 'C' */
Empty(1);
Get(1, 3); /* return nothing */
```

Ok... What are Its Unique Properties?

- For each bucket, only append-only insertion (increasing order) is allowed

```
Set (0, 1, 'A') ;  
Set (0, 2, 'B') ;  
Set (0, 3, 'C') ;  
Set (0, 0, 'D') ; /* Not Allowed */
```

- For each bucket, overwrites are prohibited

```
Set (0, 1, 'A') ;  
Set (0, 2, 'B') ;  
Set (0, 3, 'C') ;  
Set (0, 3, 'D') ; /* Not Allowed */
```

DGIBox Abstracts Modern Storage

■ Flash-based SSDs and SMR drives




Fig. 1 Conventional Track/Sector Layout




Fig. 2 Conventional Media Layout


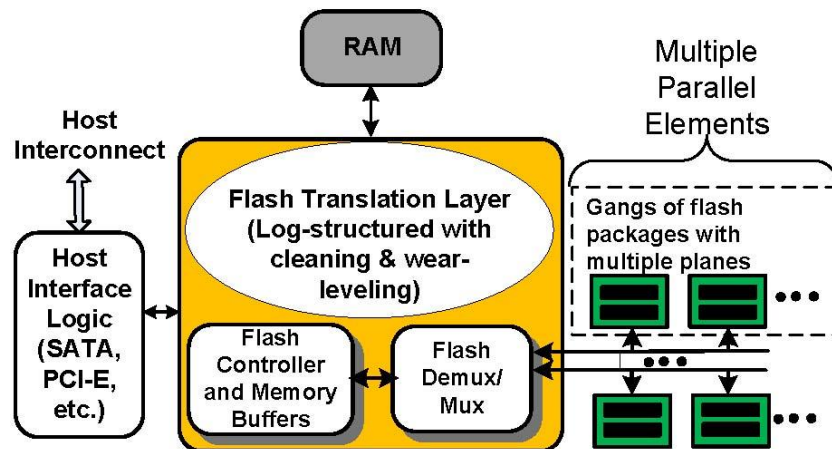



Fig. 3 Shingled Zone/Track/Sector Layout




Example

User Input:

pairs of key and value

```
write(key=0, value=A)
write(key=1, value=B)
write(key=11, value=C)
write(key=10, value=D)
write(key=9, value=E)
write(key=0, value=A')
write(key=1, value=B')
write(key=2, value=F)
write(key=3, value=G)
write(key=4, value=H)
write(key=5, value=I)
write(key=6, value=J)
```



Your Algorithm:

Map keys to buckets

```
set(0,0,A)
set(0,1,B)
set(0,2,C)
set(0,3,D)
set(1,0,E)
set(1,1,A')
set(1,2,B')
set(1,3,F)
set(2,0,G)
set(2,1,H)
set(2,2,I)
set(2,3,J)
```

Bucket₀

A
B
C
D

Bucket₁

E
A'
B'
F

Bucket₂

G
H
I
J

Example (Cont.)

User Input:
pairs of key and value

```
write(key=0, value=A)
write(key=1, value=B)
write(key=11, value=C)
write(key=10, value=D)
write(key=9, value=E)
write(key=0, value=A')
write(key=1, value=B')
write(key=2, value=F)
write(key=3, value=G)
write(key=4, value=H)
write(key=5, value=I)
write(key=6, value=J)
read(key=0)
read(key=1)
```



Your Algorithm:
Map keys to buckets

```
set(0,0,A)
set(0,1,B)
set(0,2,C)
set(0,3,D)
set(1,0,E)
set(1,1,A')
set(1,2,B')
set(1,3,F)
set(2,0,G)
set(2,1,H)
set(2,2,I)
set(2,3,J)
get(??,??)
get(??,??)
```

Dictionary
(in your algorithm)

0	(1,1)
1	(1,2)
2	(1,3)
3	(2,0)
4	(2,1)
5	(2,2)
6	(2,3)
7	
8	
9	(1,0)
10	(0,3)
11	(0,2)

Bucket₀

A
B
C
D

Bucket₁

E
A'
B'
F


Bucket₂

G
H
I
J

Example (Cont.)

User Input:
pairs of key and value

```
write(key=0, value=A)
write(key=1, value=B)
write(key=11, value=C)
write(key=10, value=D)
write(key=9, value=E)
write(key=0, value=A')
write(key=1, value=B')
write(key=2, value=F)
write(key=3, value=G)
write(key=4, value=H)
write(key=5, value=I)
write(key=6, value=J)
read(key=0)
read(key=1)
```



Your Algorithm:
Map keys to buckets

```
set(0, 0, A)
set(0, 1, B)
set(0, 2, C)
set(0, 3, D)
set(1, 0, E)
set(1, 1, A')
set(1, 2, B')
set(1, 3, F)
set(2, 0, G)
set(2, 1, H)
set(2, 2, I)
set(2, 3, J)
get(1, 1)
get(1, 2)
```

Dictionary
(in your algorithm)

0	(1,1)	
1	(1,2)	
2	(1,3)	
3	(2,0)	
4	(2,1)	
5	(2,2)	
6	(2,3)	
7		
8		
9	(1,0)	
10	(0,3)	
11	(0,2)	

Bucket₀

A
B
C
D

Bucket₁

E
A'
B'
F

Bucket₂

G
H
I
J

You can get locations of keys 0 and 1
by referring to the dictionary

Example (Cont.)

User Input:
pairs of key and value

```
write(key=0, value=A)
write(key=1, value=B)
write(key=11, value=C)
write(key=10, value=D)
write(key=9, value=E)
write(key=0, value=A')
write(key=1, value=B')
write(key=2, value=F)
write(key=3, value=G)
write(key=4, value=H)
write(key=5, value=I)
write(key=6, value=J)
read(key=0)
read(key=1)
write(key=7, value=K)
write(key=8, value=L)
```



Your Algorithm:
Map keys to buckets

```
set(0,0,A)
set(0,1,B)
set(0,2,C)
set(0,3,D)
set(1,0,E)
set(1,1,A')
set(1,2,B')
set(1,3,F)
set(2,0,G)
set(2,1,H)
set(2,2,I)
set(2,3,J)
get(1,1)
get(1,2)
no free
space!!!
```

Dictionary
(in your algorithm)

0	(1,1)
1	(1,2)
2	(1,3)
3	(2,0)
4	(2,1)
5	(2,2)
6	(2,3)
7	
8	
9	(1,0)
10	(0,3)
11	(0,2)

Bucket₀

A
B
C
D

Bucket₁

E
A'
B'
F

Bucket₂

G
H
I
J

Example (Cont.)

We actually have enough space to store values of 7 and 8
 Bucket₀ contains old data that are not valid now

User Input:
 pairs of key and value

```

write(key=0, value=A)
write(key=1, value=B)
write(key=11, value=C)
write(key=10, value=D)
write(key=9, value=E)
write(key=0, value=A')
write(key=1, value=B')
write(key=2, value=F)
write(key=3, value=G)
write(key=4, value=H)
write(key=5, value=I)
write(key=6, value=J)
read(key=0)
read(key=1)
write(key=7, value=K)
write(key=8, value=L)
  
```



Your Algorithm:
 Map keys to buckets

```

set(0,0,A)
set(0,1,B)
set(0,2,C)
set(0,3,D)
set(1,0,E)
set(1,1,A')
set(1,2,B')
set(1,3,F)
set(2,0,G)
set(2,1,H)
set(2,2,I)
set(2,3,J)
get(1,1)
get(1,2)
no free
space!!!
  
```

Dictionary
 (in your algorithm)

0	(1,1)
1	(1,2)
2	(1,3)
3	(2,0)
4	(2,1)
5	(2,2)
6	(2,3)
7	
8	
9	(1,0)
10	(0,3)
11	(0,2)

Bucket₀

	A	
	B	
	C	
	D	

Bucket₁

E
A'
B'
F

Bucket₂

G
H
I
J

Example (Cont.) – Cleaning

■ Cleaning the bucket₀

Your Algorithm:
Map keys to buckets

```
get(0, 2)  
get(0, 3)
```

```
empty(0)
```

```
set(0, 0)  
set(0, 1)
```

Five extra operations:
2 gets + 2 sets + 1 empty



These are overheads!!!

Read to memory
temporarily


Empty
Bucket₀

Write back
C and D


Bucket₀



Bucket₀



Bucket₀



Example (Cont.)

User Input:
pairs of key and value

```
write(key=0, value=A)
write(key=1, value=B)
write(key=11, value=C)
write(key=10, value=D)
write(key=9, value=E)
write(key=0, value=A')
write(key=1, value=B')
write(key=2, value=F)
write(key=3, value=G)
write(key=4, value=H)
write(key=5, value=I)
write(key=6, value=J)
read(key=0)
read(key=1)
write(key=7, value=K)
write(key=8, value=L)
```



Your Algorithm:
Map keys to buckets

```
set(0,0,A)
set(0,1,B)
set(0,2,C)
set(0,3,D)
set(1,0,E)
set(1,1,A')
set(1,2,B')
set(1,3,F)
set(2,0,G)
set(2,1,H)
set(2,2,I)
set(2,3,J)
get(1,1)
get(1,2)
no free
space!!!
```

Dictionary
(in your algorithm)

0	(1,1)
1	(1,2)
2	(1,3)
3	(2,0)
4	(2,1)
5	(2,2)
6	(2,3)
7	
8	
9	(1,0)
10	(0,1)
11	(0,0)

Bucket₀

C
D

Bucket₁

E
A'
B'
F

Bucket₂


G
H
I
J

Example (Cont.)

User Input:

pairs of key and value

```
write(key=0, value=A)
write(key=1, value=B)
write(key=11, value=C)
write(key=10, value=D)
write(key=9, value=E)
write(key=0, value=A')
write(key=1, value=B')
write(key=2, value=F)
write(key=3, value=G)
write(key=4, value=H)
write(key=5, value=I)
write(key=6, value=J)
read(key=0)
read(key=1)
write(key=7, value=K)
write(key=8, value=L)
```



Your Algorithm:

Map keys to buckets

```
set(0,0,A)
set(0,1,B)
set(0,2,C)
set(0,3,D)
set(1,0,E)
set(1,1,A')
set(1,2,B')
set(1,3,F)
set(2,0,G)
set(2,1,H)
set(2,2,I)
set(2,3,J)
get(1,1)
get(1,2)
no free
space!!!
```

Dictionary

(in your algorithm)

0	(1,1)
1	(1,2)
2	(1,3)
3	(2,0)
4	(2,1)
5	(2,2)
6	(2,3)
7	(0,2)
8	(0,3)
9	(1,0)
10	(0,1)
11	(0,0)

Bucket₀

C
D
K
L

Bucket₁

E
A'
B'
F

Bucket₂

G
H
I
J

Example (Cont.) – Summary

■ Operations that are performed by your algorithm

- **Input ops:** 14 writes and 2 reads
- **Extra ops:** 2 writes + 2 reads + 1 empty
- **Total:** 16 writes, 4 reads, and 1 empty

■ Analysis

- **The cost of a write:** $(16 \text{ writes} + 2 \text{ reads} + 1 \text{ empty}) / (14 \text{ writes}) = \Theta(1.36)$
 - Assume that the costs of a write, a read, and an empty are the same
- **The cost of a read:** $1 \text{ read} = \Theta(1)$

What is the Problem with the Example

■ The previous solution requires lots of DRAM

- Suppose that DGIBox has 2^{23} buckets whose sizes are 256 and the size of a table entry is 4 bytes
- The number of entries in the dictionary is $2^{23} \times 2^9 = 2^{32}$
- The amount of DRAM required for the dictionary is

$$2^{32} \times 4 \text{ bytes} = \mathbf{16 \text{ GB!!!}}$$

■ How to reduce the memory requirement?

- Hash?

Example with Hash

User Input:
pairs of key and value

```
write(key=0, value=A)
write(key=1, value=B)
write(key=11, value=C)
write(key=10, value=D)
write(key=9, value=E)
write(key=0, value=A')
write(key=1, value=B')
write(key=2, value=F)
write(key=3, value=G)
write(key=4, value=H)
write(key=5, value=I)
write(key=6, value=J)
read(key=0)
read(key=1)
write(key=7, value=K)
write(key=8, value=L)
```


Hash Function:
Division

```
0/4=0, 0%4=0
1/4=0, 1%4=1
11/4=2, 11%4=3
10/4=2, 10%4=2
```

Your Algorithm:
Map keys to buckets

```
set(0, 0, A)
set(0, 1, B)
set(2, 3, C)
set(2, 2, D) ✱
```

Bucket₀



Bucket₁



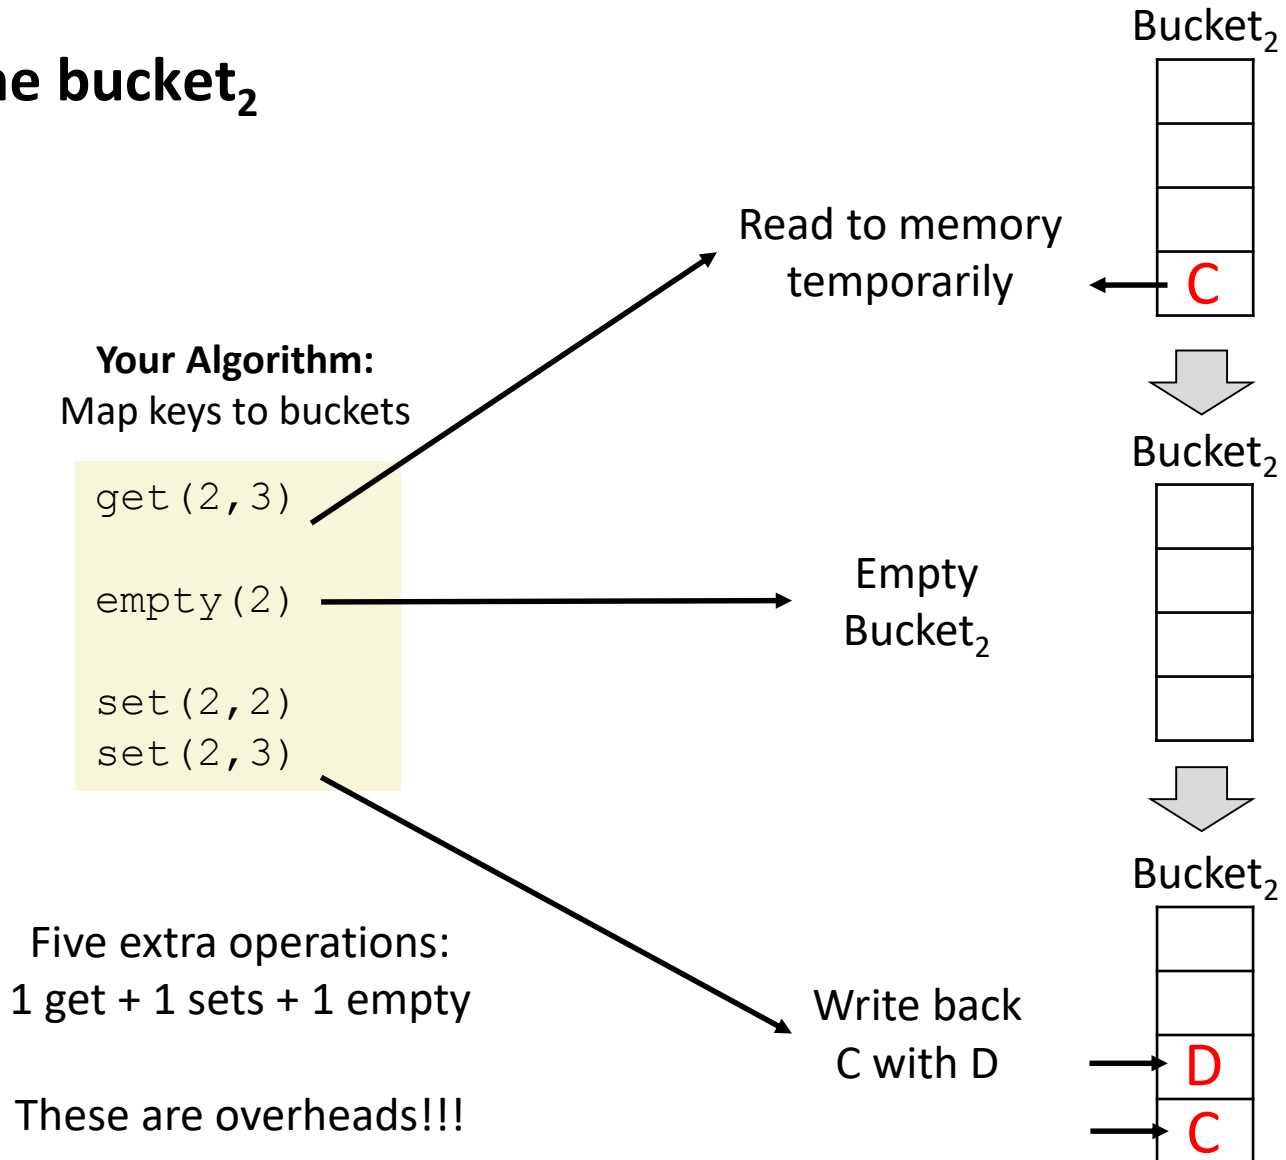
Bucket₂



It violates the property of DGIBox ✱ D

Example with Hash (Cont.)

■ Cleaning the bucket₂



Example with Hash (Cont.)

User Input:
pairs of key and value

```
write(key=0, value=A)
write(key=1, value=B)
write(key=11, value=C)
write(key=10, value=D)
write(key=9, value=E)
write(key=0, value=A')
write(key=1, value=B')
write(key=2, value=F)
write(key=3, value=G)
write(key=4, value=H)
write(key=5, value=I)
write(key=6, value=J)
read(key=0)
read(key=1)
write(key=7, value=K)
write(key=8, value=L)
```

Hash Function:
Division

```
0/4=0, 0%4=0
1/4=0, 1%4=1
11/4=2, 11%4=3
10/4=2, 10%4=2
9/4=2, 9%4=1
0/4=0, 0%4=0
1/4=0, 1%4=1
2/4=0, 2%4=2
3/4=0, 3%4=3
4/4=1, 4%4=0
5/4=1, 5%4=1
6/4=1, 6%4=2
0/4=0, 0%4=0
1/4=0, 1%4=1
7/4=1, 7%4=3
8/4=2, 8%4=0
```

Your Algorithm:
Map keys to buckets

```
set(0,0,A)
set(0,1,B)
set(2,3,C)
set(2,2,D)
set(2,1,E)
set(0,0,A')
set(0,1,B')
set(0,2,F)
set(0,3,G)
set(1,0,H)
set(1,1,I)
set(1,2,J)
get(0,0)
get(0,1)
set(1,3,K)
set(2,0,L)
```

Bucket₀

A'
B'
F
G

Bucket₁

H
I
J
K

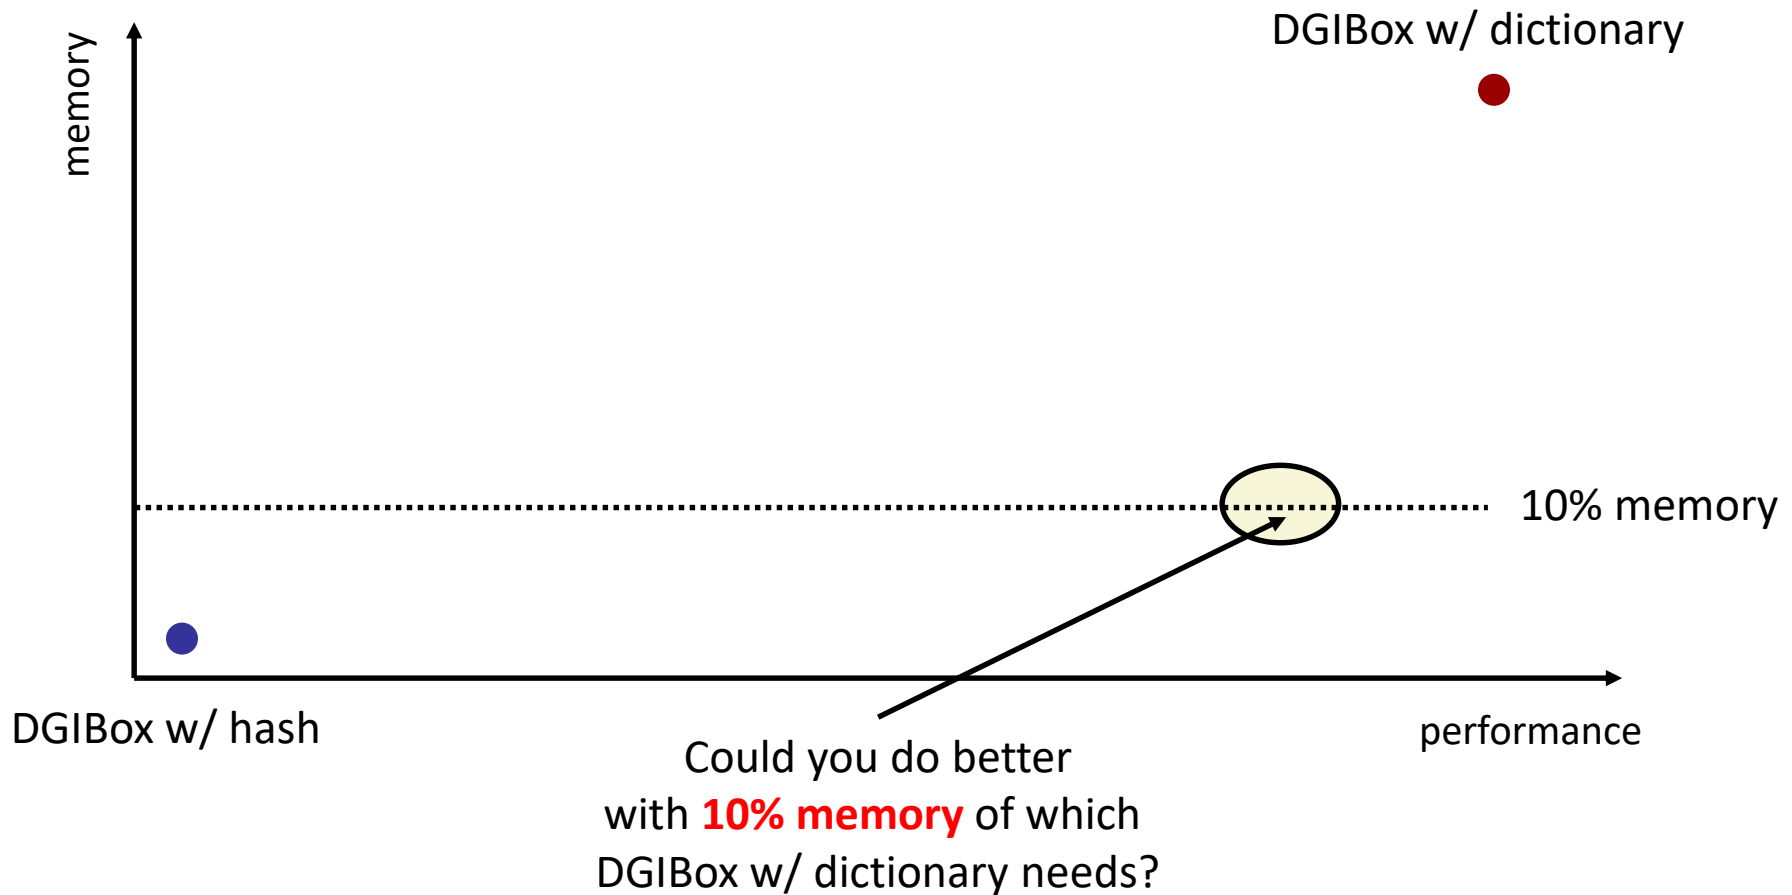
Bucket₂

L
E
D
C

Example (Cont.) – Summary

- DGIBox with Hash requires zero memory for the dictionary,
- But, operations that are performed
 - Input ops: 14 writes and 2 reads
 - Extra ops: 8 writes + 8 reads + 5 empty
 - Total: 22 writes, 10 reads, and 5 empty
- Analysis
 - The cost of a write: $(22 \text{ writes} + 10 \text{ reads} + 5 \text{ empty}) / (14 \text{ writes}) = \Theta(2.64)$
 - Assume that the costs of a write, a read, and an empty are the same
 - The cost of a read: $1 \text{ read} = \Theta(1)$

Trade-off Between Cost and Memory



Grading

- **Develop algorithms providing basic operations (write and read) that operate on DGIBox**

- **Specific Implementation Topics**
 - 1. DGIBox w/ Hash: **10%**
 - 2. DGIBox w/ dictionary: **20%**
 - 3. DGIBox w/ your own algorithm: **40%**
 - 1. Correctly works with limited memory (10% of dictionary): **10%**
 - 2. Performance: **30%**
 - It is a race!
 - 4. Design Report: **30%**
 - 1. It would be good if you formulate complexity of your algorithm

Rules

■ Development Tool & Languages

- Language: C/C++ (No Java and Python!)
- Microsoft Visual Studio

■ What to submit?

- Source codes that include all your algorithms
- Design report

■ Where?

- To TA via email

■ Until when?

- Until 11:59pm on November 15th
 - Double column, **3 pages**, 10 font size, No title page
 - Template & Sample: Uploaded in BlackBoard

Cheating?

- **F Grade**
- **TA will check up all your codes manually and using cheating detection tools**